Inverting Python: Designing NTP-Prior-Resistant Programming Languages toward LLM-Resistant Coding Assessment

Jaeyeong CHOI
DGIST EECS
jaeyeong2022@dgist.ac.kr

***Abstract*—LLMs now solve standard programming tasks well enough to undermine their value as student assessments. We investigate whether a human-interpretable programming language can defeat current-generation LLM next-token-prediction (NTP) priors. We propose a four-level taxonomy—L1 (token alias), L2 (syntax inversion), L3 (explicit semantic inversion), L4 (implicit semantic inversion)—and evaluate eleven models: seven OpenAI core models (`gpt-4o-o4-mini`), frontier models `gpt-5.4/o3`, and two open-weight models (`Llama-3.3-70B`, `Qwen3-32B`) via Groq. L1–L3 provide partial or capability-dependent resistance; L4 defeats *all* tested models on prior-entrenched tasks through *pattern blindness*—including frontier models `gpt-5.4` and `o3` (0/10 each, PPR=1.0 on Fibonacci). A multi-task extension (345 runs; Fibonacci, `is_sorted`, `count_vowels`) reveals task-dependent resistance and a third failure mode, *operational substitution*. An annotation density experiment (E34, 150 runs) yields the paper's most novel finding: *pattern mimicry via code examples outperforms verbal rule annotation at every density level*, with code-example conditions achieving 50–80% pass while fully-annotated conditions collapse to 10–30%. A redesigned I/O transparency experiment (E35, 60 runs) establishes I/O opacity as a second, complementary resistance pathway: opaque I/O pairs alone produce universal resistance (0/30, 0%), independent of prior strength. Across 1150 total L4 evaluations (11 models, 9 OpenAI + 2 open-weight), we identify three failure mechanisms: *prior dominance*, *pattern blindness*, and *operational substitution*.[1] A chain-of-thought ablation (160 runs, 4 models) reveals a *reasoning–generation dissociation*: models articulate the inversion rule when prompted to reason yet still generate standard Python code. Open-weight evaluation confirms L4 pattern blindness with distinct failure modes and wide L1 compliance variance (KLR=0.23 vs. 1.00), suggesting architecture and fine-tuning govern alias resistance.**

***Index Terms*—confusion language, LLM resistance, NTP prior, semantic inversion, programming education, code assessment**

[1]Count breakdown: 350 (strict ablation, 7 OpenAI core, $n = 50$) + 75 (open-weight ablation: Llama $n_{\text{eff}} = 25$ + Qwen $n = 50$) + 20 (E24 frontier: gpt-5.4 + o3, $n = 10$ each) + 120 (hard-task, 4 models $\times$ 3 tasks $\times$ $n = 10$) + 345 (multi-task, 7 models) + 40 (E28 annotation density gradient, `o4-mini`, $n = 10\times4$ levels) + 150 (E34 annotation density curve, 3 models $\times$ 5 levels $\times$ $n = 10$) + 50 (E32 cross-task transfer: `gpt-4o` $n = 30$ + `gpt-4.1-mini` $n = 10$ + `o4-mini` $n = 10$). CoT ablation (160 runs) reported separately to avoid double-counting Variant A runs. E32-ctrl control condition (10 runs, `gpt-4o`) and E33/E33v2 I/O transparency (120 runs total) excluded as confounded; E35 (60 runs) reported separately as a complementary resistance pathway.

I. Introduction

Modern code-generation benchmarks [1], [2] show that large language models (LLMs) can solve standard programming tasks with high accuracy. This creates a critical challenge for programming education: if LLMs solve assignments faster and more accurately than students, how can instructors meaningfully assess student understanding?

This paper studies a complementary but understudied question:

> *Is it possible to design a programming language that is human-interpretable but systematically defeats LLM next-token-prediction priors—making it resistant to AI-based code generation?*

The core asymmetry. The key insight is that humans and LLMs fail in fundamentally different ways when given an unfamiliar language rule. A human student who reads the rule *"if-blocks execute when the condition is **false**"* can immediately apply it—the rule is explicit, verifiable, and overrides prior habit. An LLM given the identical rule in its prompt instead produces standard Python: its next-token prediction (NTP) prior, built from billions of Python training tokens, systematically overrides the stated instruction. This asymmetry—*humans follow novel rules; LLMs revert to training priors*—is the foundation of our resistance framework and requires no additional experiment to accept: it follows directly from the mechanics of NTP-based generation.

We decompose this into three sub-questions:

1) Does changing token identifiers (L1: alias substitution) create effective LLM resistance?
2) Does inverting syntactic structure (L2) create effective resistance?
3) Does inverting semantic behavior (L3/L4) create deeper resistance?

Our key empirical findings are:

- **L1 (token substitution):** Partially effective—keyword leakage rate (KLR) ranges from 0.26 to 0.69 across models. Models revert to Python priors at measurable rates.
- **L2 (syntax inversion):** Not effective—models adopt new syntax completely (SIR=0.0) when given explicit rules.
- **L3 (semantic inversion, explicit):** Capacity-dependent—weaker models fail completely (PPR=1.0

for `gpt-4o-mini`), while stronger models succeed (PPR=0.0 for `gpt-4o` and `gpt-4.1-mini`).

- **L4 (semantic inversion, implicit):** Effective against all tested models on prior-entrenched tasks (e.g., Fibonacci)—even strong models exhibit *pattern blindness*. However, capable models (gpt-5.4, o3) succeed on complex algorithmic tasks (merge sort, BFS) where prior entrenchment is shallow.

A key boundary condition: L4 resistance requires tasks with deep, uniform Python priors. We identify *prior entrenchment depth* as a third design axis beyond perturbation level and rule delivery mode.

We also identify three distinct LLM failure modes: *prior dominance*, where the prior overrides stated rules; *pattern blindness*, where the model fails to extract semantic rules from examples; and *operational substitution*, where the model achieves correct output via arithmetic workaround rather than adopting the semantic rule.[2]

This framework provides both a theoretical taxonomy for LLM-resistant language design and practical tools for constructing programming assessments that remain meaningful in the LLM era.

II. Related Work

A. Code Generation Benchmarks and Evaluation

HumanEval [1] and MBPP [2] established functional correctness as the standard metric for code-generation evaluation. More recent benchmarks scale and diversify tasks: BigCodeBench [3] adds complex, multi-function instructions, and LiveCodeBench [4] avoids contamination through continuous updating. Across all these benchmarks, Python dominates training corpora, embedding strong token-level priors into all models evaluated. This motivates our L1 baseline: even with explicit alias instructions, models leak Python keywords at 26–69%.

B. Code-Oriented LLM Architectures and Training

Pre-trained code models—CodeBERT [5], CodeT5 [6], and InCoder [7]—demonstrate that NTP objectives on Python code reinforce keyword position priors. More recent instruction-tuned models (WizardCoder [8], DeepSeek-Coder [9]) amplify these priors through even larger code corpora. This trajectory predicts our finding: *stronger code LLMs do not necessarily mean better alias compliance*—gpt-5.4-mini (KLR=0.67) underperforms gpt-4o (KLR=0.37).

C. Instruction-Following in Code (2024–2025)

CodeIF [10] and MultiCodeIF [11] introduce benchmarks measuring LLM adherence to developer-specified constraints in code generation tasks. Critically, MultiCodeIF reports that models perform relatively well on *explicit, concrete* constraints but degrade substantially on *implicit or abstract* constraints, and that performance drops sharply as the number of hierarchical constraints increases. This explicit–implicit asymmetry is the closest mechanistic evidence in the literature to our L3/L4 findings: models can follow explicit semantic inversion rules (L3, with some model-capacity dependence), yet fail entirely when the same semantic inversion is delivered only through examples without explicit articulation (L4, pattern blindness across all tested models). The key distinction is scope: CodeIF and MultiCodeIF benchmark *constraint-following* within standard programming tasks, whereas our work *designs* a programming language whose semantic structure systematically conflicts with NTP priors, making resistance a structural property rather than a task-level benchmark artifact.

D. Instruction Hierarchy and Prior Conflicts

Chen et al. [12] propose an instruction hierarchy to resolve conflicts between system-level and user-level instructions. Their framework addresses natural-language directive conflicts; ours involves programming-language semantic rule conflicts—a dimension not studied in alignment work. Xie et al. [13] and Wu et al. [14] show that when context conflicts with parametric knowledge, LLMs often default to pretraining priors. Zhang et al. [15] survey this knowledge conflict problem broadly. Shi et al. [16] demonstrate LLMs are easily distracted by irrelevant context; our L4 condition represents the inverse: models *fail to use* relevant semantic cues provided in examples.

E. NTP Priors and In-Context Learning

Wang et al. [17] formalize a "law of next-token prediction" providing theoretical grounding for how NTP training shapes output distributions. Min et al. [18] show ICL demonstrations affect outputs primarily through distributional surface patterns rather than semantic content—explaining why models fail to extract inverted semantics from L4 examples. Dong et al. [19] survey ICL comprehensively. Lampinen et al. [20] document "A-Not-B" errors where LLMs fail to suppress previously established response patterns under context changes, directly analogous to our L4 pattern blindness.

F. Counterfactual and Semantic Reasoning Failures

Zhao et al. [21] show LLMs struggle when counterfactual reasoning requires overriding parametric knowledge—the same mechanism we observe at L3 for weak models. Turpin et al. [22] show LLM self-explanations are often unfaithful to actual behavior. Valmeekam et al. [23] document reasoning rigidity—models revert to memorized strategies under redefined rules—providing further precedent for our pattern blindness finding.

G. Adversarial Perturbation and Code Resistance

PromptBench [24] demonstrates LLM fragility under text perturbations. Pan et al. [25] show semantically aligned code prompt perturbations cause up to 42% accuracy drops. Jha et al. [26] use structured adversarial prompts to break LLM reasoning in code tasks. Khan et al. [27] specifically design prompt perturbations to reduce LLM correctness on programming assignments by 77%, most directly related to our goal. Our approach is complementary: instead of perturbing

[2]Operational substitution was observed on `count_vowels` for `gpt-4o` at $n = 50$.

prompts, we design a confusion language whose NTP prior conflict is a *structural property*, making it principled and reproducible.

H. DSL and Constrained Code Generation

Wu et al. [28] survey LLM performance on low-resource and domain-specific languages, confirming that Python-trained models underperform on less common languages. ConCodeEval [29] benchmarks code constraint comprehension in DSLs; L2CEval [30] evaluates language-to-code generation across specialized languages. Constrained decoding [31], [32] enforces grammar-level output validity. Our L2 experiment shows models adopt explicit syntax rules entirely, but semantic constraints (L3/L4) prove far more resistant—showing grammar-level constraints are insufficient without semantic-level challenges.

I. Real-World Engineering Benchmarks

SWE-bench [33] reveals that LLMs fail to generalize from canonical benchmarks to real repository-level tasks—confirming that high performance on standard benchmarks does not indicate robustness to structural novelty, as we observe in our L1 L4 evaluations.

J. LLM in Programming Education

Dakhel et al. [34] show LLM coding tools are an "asset or liability" depending on student expertise. Lau and Guo [35] document novice programmers adopting LLM outputs uncritically. Prather et al. [36] show mixed student perceptions. Liffiton et al. [37] propose guardrailed LLM tutoring. These studies confirm a shared concern: LLM integration can undermine independent skill development.

K. Academic Integrity and LLM-Resistant Assessment

Novak and McDanel [38] propose contextual, personalized assignments to resist LLM cheating—the most directly related paper to our educational goal—but without a formal theory of what features drive LLM failure. Code perplexity detection [39] is brittle; our taxonomy offers a principled design-based solution. RLHF-based alignment [40] improves instruction-following but, as our L3/L4 results show, remains insufficient when semantic inversions must be extracted from examples.

L. Research Gap

Prior adversarial work degrades generation quality without principled language design. Prior DSL work measures performance without explaining failure mechanisms. Prior knowledge-conflict work focuses on factual knowledge, not programming-language semantics. Our contribution closes these gaps: a theory-grounded L1–L4 taxonomy that characterizes LLM failure by mechanism (prior dominance vs. pattern blindness) and guides principled LLM-resistant language design.

TABLE I: Four-level confusion language taxonomy. Each level targets a distinct linguistic dimension and delivery modality. L4 is the only level where all tested models fail regardless of capacity.

Level	Perturbation	Delivery	Observed failure mode
L1	Token alias	Explicit	Prior dominance (partial)
L2	Syntax inv.	Explicit	None—adopted
L3	Semantic inv.	Explicit	Capacity-dependent
L4	Semantic inv.	Implicit	Pattern blindness

Note: L4 experiments use three distinct sub-conditions across sections: (1) strict example-only ablation (§V-D, $n = 50$), the canonical L4 condition; (2) partially-annotated multi-task (§V-E), a weakened form with inline comment hints; and (3) hard-task extension (§V-N), example-only on algorithmically complex tasks. Findings differ by sub-condition; see §VI-C for synthesis.

III. Method

A. Research Questions and Hypotheses

RQ1: At what language design level does LLM resistance become robust? **RQ2**: What failure mechanism is active at each level—prior dominance or pattern blindness?

H1: Token-level alias substitution (L1) creates partial but not complete resistance. **H2**: Syntactic inversion (L2) is insufficient—models adopt new syntax when given explicit rules. **H3**: Semantic inversion (L3/L4) provides robust resistance, but differs in mechanism by model capacity.

B. Language Level Taxonomy

We design a structured family of programming language variants with properties engineered to conflict with LLM next-token prediction priors—which we term a *confusion language* (our proposed terminology). We distinguish confusion languages from esoteric languages (e.g., Brainfuck), which aim for human uninterpretability; confusion languages retain full human interpretability while selectively targeting LLM NTP priors. Each confusion language variant defines rules along two axes: (1) *perturbation depth* (token, syntax, semantic) and (2) *rule delivery* (explicit vs. implicit). Table I summarizes the resulting four-level taxonomy.

Figure 1 illustrates the two-axis design space and the location of each level. The two gray cells (Token/Implicit and Syntax/Implicit) are theoretically possible configurations that we did not study; Token/Implicit would expect higher KLR than L1-Explicit (weaker contextual cues), while Syntax/Implicit would expect higher SIR. We leave these as open positions in the design space for future work.

C. L1: Token Alias Substitution

Python keywords (`if`, `for`, `in`, `return`, `def`) are replaced with aliases. Constraints: one-to-one mapping, no alias collision, token-boundary replacement, literal/comment protection.

D. L2: Syntax Structure Inversion

Function definitions, conditionals, loops, and block delimiters are structurally inverted relative to Python. Example: `def fib(n):` becomes `:define fib [n] ->`. Model is given explicit grammar rules.

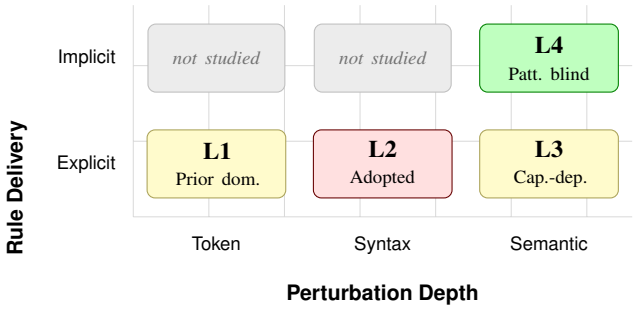

Fig. 1: Two-axis confusion language design space. Perturbation depth (token/syntax/semantic) and rule delivery (explicit/implicit) together determine LLM failure mode. L4 (green = most resistant) occupies the hardest cell: semantic perturbation without explicit rule—the only condition where all tested models fail.

Python (standard)

```
def fib(n):
    if n <= 1:
        return n
    return fib(n-1) + fib(n-2)
```

Standard semantics: if-block runs when condition is TRUE

L3 (semantic inversion)

```
def fib(n):
    if n <= 1:
        return fib(n-1) + fib(n-2)
    return n
```

Inverted: if-block runs when condition is FALSE

Fig. 2: Semantic inversion in L3. Surface syntax is identical to Python, but if-block semantics are inverted. The correct L3 solution (right) looks counter-intuitive to Python-trained models, which generate the left version regardless of the stated rule.

E. L3: Semantic Inversion (Explicit Rules)

Python surface syntax is retained but the semantics of if-blocks are inverted: the indented block executes when the condition is *false*. Model is explicitly told this rule. Figure 2 illustrates the concept.

F. L4: Semantic Inversion (Implicit, Example-Only)

Same as L3, but no rule is stated. Instead, models receive 2–3 worked examples that demonstrate the inversion. Models must extract the semantic pattern through few-shot reasoning.

G. Evaluation Metrics

Redundant metrics are removed. We report two independent measures:

- **KLR** (Keyword Leakage Rate): fraction of alias slots where original Python keyword was used. Note: ACR $= 1 -$ KLR and PIR $\approx$ KLR are derived quantities and not separately reported.
- **PSS** (Partial Structural Score, exploratory): weighted sub-component credit (alias: 0.4, parse: 0.25, skeleton: 0.25, format: 0.1). Weights are heuristic; PSS is reported as a secondary diagnostic alongside KLR and should not be used as the sole basis for compliance conclusions.
- **PPR** (Python Prior Rate, L3/L4): fraction of samples using Python-style condition (no inversion).
- **PES** (Prior Entrenchment Score, exploratory): probability assigned by the model to the canonical Python token at the key decision point of a task's base case, measured via the logprobs API (`top_logprobs=5`, `max_tokens=1`, `temperature=0`). PES quantifies the strength of the prior that L4 must overcome; see E31 results in §V-J.

H. Threats to Validity

Internal: judge rules may be stricter than human acceptance; PSS weights are heuristic. **Construct**: multi-task extension (T1 fib, T2 is_sorted, T3 count_vowels) partially addresses single-task limitation; broader coverage (nested control-flow, stateful logic) remains for future work. Open-weight evaluation (`Llama-3.3-70B`, `Qwen3-32B`) is reported with extended 5-variant runs (§V-H): `Qwen3-32B` completed 50/50 valid runs (full replication); `Llama-3.3-70B` achieved 25/50 valid runs ($n_{\text{eff}} = 5$ per variant due to Groq rate limiting), insufficient for per-variant statistical claims. **External**: primary results cover nine model families (seven OpenAI + two open-weight via Groq); `Llama-3.3-70B` per-variant results remain underpowered and full replication is deferred to future work. Human baseline (Exp-5) needed to validate the educational claim.

IV. Experiments

A. Experimental Overview

We ran experiments at each language level for available model families. All runs used prompts from `v1-v120` (L1) or dedicated level-specific banks (L2–L4). API authentication was managed via project environment variables.

B. Models Evaluated

- **Core OpenAI** (seven models): `gpt-4o`, `gpt-4o-mini`, `gpt-4.1`, `gpt-4.1-mini`, `gpt-4.1-nano`, `gpt-5.4-mini`, `o4-mini`
- **Frontier extension**: `gpt-5.4`, `o3` (hard-task extension §V-N and E24 Fibonacci ablation §V-D)
- **Open-weight via Groq**: `Llama-3.3-70B` ($n_{\text{eff}} = 25$, per-variant underpowered), `Qwen3-32B` ($n = 50$, full replication)

C. L1 Experiment: Token Alias Substitution

120 prompts per model, versioned v1–v120 with varied alias mappings. Context-pack ablation: baseline vs. explicit alias contract system prompt.

D. L2 Experiment: Syntax Inversion

5 prompts $\times$ 4 runs each ($n = 20$) using `:define ... ->` syntax (see Appendix). Model given complete grammar specification. Evaluated for `gpt-4o`, `gpt-4.1-mini`, `gpt-4.1-nano` (plus $n = 5$ for `gpt-4o-mini` from earlier run).

E. L3 Experiment: Semantic Inversion (Explicit)

10 prompts $\times$ 2 runs each ($n = 20$). System message explicitly states: "if-blocks execute when condition is FALSE." Evaluated for `gpt-4o`, `gpt-4.1-mini`, `gpt-4.1-nano` (plus $n = 5$–10 for `gpt-4o-mini` and `gpt-4o` from earlier runs).

F. L4 Experiment: Semantic Inversion (Implicit)

Initial screening: 10 prompts per model, no rule stated; only 2 worked examples with inverted semantics shown. Model must infer pattern. Ablation (n=50): 5 prompt variants $\times$ 10 runs each, varying the number of worked examples (2 vs. 3), presence of inline comments, order of examples, and domain of examples. Evaluated across all seven models at $n = 50$ each (350 total L4 ablation runs): `gpt-4o`, `gpt-4o-mini`, `gpt-4.1`, `gpt-4.1-mini`, `gpt-4.1-nano`, `gpt-5.4-mini`, and `o4-mini`.

G. L4 Multi-task Generalization Experiment

To assess whether pattern blindness generalizes across algorithmic domains, we extended the L4 evaluation to two additional tasks beyond Fibonacci: **T2 (is_sorted)**: write a function returning True if a list is ascending, and **T3 (count_vowels)**: count vowel characters in a string. Both use a shared prompt header with two worked examples demonstrating inverted semantics plus inline code comments that *partially* annotate the inversion rule (e.g., `# runs when condition is FALSE, i.e., ...`). This is a *partially-annotated L4* condition—more informative than the strict example-only Variant A ablation (§V-D) and therefore **not directly comparable** to Table IV. We ran $n = 10$ per task for `gpt-4o`, `gpt-4o-mini`, `gpt-4.1`, `gpt-4.1-nano`, and `gpt-5.4-mini`; $n = 20$ for `gpt-4.1-mini`; and $n = 15$ for `o4-mini` (255 initial runs). We then scaled `gpt-4o` T3 and `gpt-4.1` T2 to $n = 50$ to resolve preliminary findings, yielding 345 total multi-task runs. Results are in Table V.

H. CoT Ablation Experiment

To test whether chain-of-thought reasoning overcomes pattern blindness, we added a CoT instruction ("think step by step about the pattern first, then write code") to the L4 Variant A prompt. We compared CoT vs. no-CoT for four models: `gpt-4o`, `gpt-4.1-mini`, `gpt-4.1-nano`, and `o4-mini` ($n = 20$ per condition, 160 total runs). Judge criteria are identical to the L4 ablation; we additionally track whether the model's reasoning text mentions the inversion pattern (operationally defined as containing at least one of the substrings: `"inverted"`, `"opposite"`, `"false when"`, or `"condition is false"`, case-insensitive regex match on the full reasoning output).

I. L4 Hard-Task Extension: Complex Algorithms and Latest Models

To test whether task complexity amplifies or attenuates pattern blindness, we designed three "hard" L4 tasks requiring multi-step algorithms with deep Python priors: **H1 (merge_sort)**: implement merge sort that returns *descending* order (prior: ascending); **H2 (binary_search)**: return -1 if element is *found*, insertion index if *not found* (prior: index if found, -1 if not); **H3 (bfs)**: return BFS traversal in *reverse* order (prior: standard BFS order). Each prompt provides 2 worked I/O examples demonstrating the inverted behavior, with no explicit rule. We evaluated four models: `gpt-4o` (baseline), `gpt-5.4` (latest dense), `o3` (reasoning), and `o4-mini` (reasoning); $n = 10$ per model per task (120 total runs). Results are in Table XI.

J. E24: Frontier Model L4 Fibonacci Ablation

To directly address whether frontier models (`gpt-5.4`, `o3`) can overcome pattern blindness on prior-entrenched tasks—a gap in the hard-task extension—we ran the strict L4 Fibonacci ablation (Variant A, example-only, no explicit rule) on both models at $n = 10$ each. This is the same Variant A condition used in Table IV. Results are reported in §V-D.

K. E28: Annotation Density Gradient (o4-mini, Pilot)

Pilot for E34. To map the annotation sensitivity boundary for reasoning models—motivated by the dissociation between CoT ($1/20$ pass) and partial annotation ($8/15$ pass) for `o4-mini`—we ran a density gradient experiment varying the fraction of inline annotation comments across the 8 worked examples in a T1 Fibonacci L4 prompt. Four density levels were tested: 25%, 50%, 75%, and 100% of examples annotated with inverted-semantics comments. $n = 10$ per density level (40 total runs, `o4-mini` only). This pilot directly operationalizes the annotation-sensitivity hypothesis; the full multi-model replication is E34 (§IV-P).

L. E25–E27: Diagnostic L3 for `gpt-4.1` and `Qwen3-32B`

E25: To determine whether `gpt-4.1`'s L4 T2 failure (1/50) reflects incapacity for semantic inversion or purely pattern-extraction failure, we ran L3 T2 evaluations (`is_sorted`, explicit rule) for `gpt-4.1` at $n = 10$. E26/E27: To determine whether `Qwen3-32B`'s L4 failure is L4-specific or reflects a deeper L3 instruction-following deficit, we ran L3 Fibonacci evaluations at $n = 5$ (E26 pilot) and $n = 20$ full replication (E27). Results are in §VI-C and §VI-D.

M. E31: Prior Entrenchment Score via Logprobs

To directly quantify prior strength at the key generation decision point, we queried the logprobs API (`top_logprobs=5`, `max_tokens=1`, `temperature=0`) at the token immediately following `return` for four tasks (fib, merge_sort, is_sorted, count_vowels) across two models (`gpt-4o`, `gpt-4.1-mini`). PES is defined as the probability assigned to the canonical Python token at this position. This directly measures the prior that L4 must overcome. Results are in §V-J.

N. E32: Cross-Task Transfer

To test whether L4 resistance survives cross-task contamination, we evaluated L4 Fibonacci pass rate after showing inverted-semantics code examples from two unrelated tasks (merge sort, binary search) in the same prompt. `gpt-4.1-mini` and `o4-mini`: $n = 10$; `gpt-4o`: scaled to $n = 30$. Control condition (E32-ctrl): normal (non-inverted) sort examples preceding fib, `gpt-4o`, $n = 10$. Temperature=0.7 for all. Results are in §V-K.

O. E33 and E35: I/O Transparency Experiments

E33 tested IO-A (transparent Fibonacci-adjacent mapping) vs. IO-B (opaque counter function with disclosed algorithm), revealing a confound: IO-B disclosed algorithmic steps, making both tasks procedurally transparent. Models (`gpt-4o`, `gpt-4.1-mini`, `gpt-4o-mini`) achieved 10/10 on both. E33v2 replicated with `gpt-4o`, `gpt-4.1-mini`, `o4-mini` to confirm confound robustness. E35 (redesigned) replaced IO-B with pure I/O pairs (8 examples, no algorithm description) to cleanly isolate opacity as an independent variable. $n = 10$ per cell, temperature=0.7. Results are in §V-L.

P. E34: Annotation Density Curve (Multi-Model, Full Replication)

Full replication of the E28 pilot (§IV-K), extending to three models (`gpt-4o`, `gpt-4.1-mini`, `o4-mini`) across five density levels (0%, 25%, 50%, 75%, 100%) on L4 Fibonacci with inverted comparisons. The 0% condition includes an inverted *code* example (not I/O-only), making it qualitatively distinct from verbal annotation levels (25–100%). $n = 10$ per cell (150 total runs). Results are in §V-M.

Q. Reproducibility

All experimental artifacts, commands, and results are committed to the repository under `docs/research/results/`. See reproducibility protocol section for commands.

V. Results

A. L1: Partial resistance; model-family compliance patterns

All OpenAI models show non-zero KLR (0.21–0.69), confirming prior dominance. Three findings stand out.

(1) `gpt-5.4-mini` (KLR=0.67) underperforms `gpt-4o` (KLR=0.37) despite being a newer model; because these models differ in training corpus and RLHF regime—not just scale—we characterize this as *model-family-specific compliance patterns* rather than a scale effect. The observed difference may reflect stronger code-generation fine-tuning in `gpt-5.4-mini` (a stated hypothesis, not a measured mechanism).

(2) For models with paired conditions, context-pack delivery reduces KLR: `gpt-4o-mini` (0.42 $\rightarrow$ 0.26, -38%) and `gpt-4.1-mini` (0.43 $\rightarrow$ 0.21, -51%).

(3) The reasoning model `o4-mini` (KLR=0.27) performs comparably to `gpt-4o-mini` ctx-pack (0.26), suggesting chain-of-thought reasoning provides modest benefit for alias compliance.

TABLE II: L1 evaluation: Keyword Leakage Rate (KLR) per model. Baseline rows use $n = 120$; ctx-pack rows use $n = 20$. Models sorted by KLR (ascending). Context-pack reduces KLR for paired models: `gpt-4o-mini` (-38%), `gpt-4.1-mini` (-51%). Open-weight models (marked *) via Groq API, ctx-pack only. PSS (heuristic diagnostic) reported in Appendix X-C.

Model	Delivery	KLR↓
gpt-4.1-mini	ctx-pack	**0.21**
Llama-3.3-70B*	ctx-pack	0.23
gpt-4o-mini	ctx-pack	0.26
o4-mini	ctx-pack	0.27
gpt-4.1	ctx-pack	0.30
gpt-4o	baseline	0.37
gpt-4.1-nano	ctx-pack	0.38
gpt-4o-mini	baseline	0.42
gpt-4.1-mini	baseline	0.43
gpt-5.4-mini	baseline	0.67
gpt-4.1-nano	baseline	0.69
Qwen3-32B*	ctx-pack	1.00

*Open-weight via Groq API; $n = 20$, ctx-pack only. See §V-H.
Green : KLR$\leq$0.27; Red : KLR$\geq$0.67.

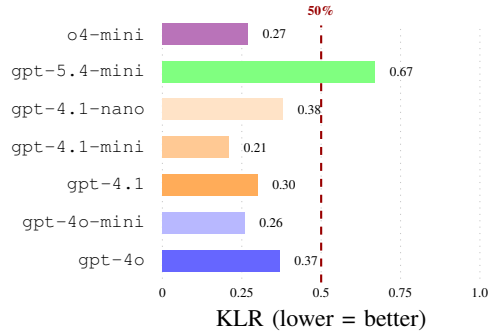

Fig. 3: L1 KLR per model (OpenAI models only). Ctx-pack condition shown for all models except `gpt-4o` and `gpt-5.4-mini` (baseline only; ctx-pack not evaluated). $n\geq20$ per bar. Open-weight models (`Llama-3.3-70B`, `Qwen3-32B`) appear in Table II only; bar chart shows OpenAI models for visual clarity. ■ `gpt-4o` fam.; ■ `gpt-4.1` fam.; ■ `gpt-5.4-mini` (baseline); ■ `o4-mini`.

Note: `gpt-4.1`, `gpt-4.1-nano`, and `o4-mini` were only evaluated under ctx-pack ($n = 20$, v1–v20); baseline comparison for these models is not available and should be interpreted with caution given the smaller and non-overlapping prompt subset.

Open-weight preliminary: `Llama-3.3-70B` (KLR=0.23) is comparable to the best OpenAI ctx-pack results; `Qwen3-32B` (KLR=1.00) is the worst result across all tested models. This cross-family variance (0.23 vs. 1.00) exceeds the spread within the `gpt-4.x` family, indicating that architecture and instruction-following fine-tuning—not parameter count—govern alias compliance. See §V-H for full details and limitations.

B. L2: Syntax inversion not resistant

Three models (`gpt-4o`, `gpt-4.1-mini`, `gpt-4.1-nano`) adopted L2 syntax fully at $n = 20$

TABLE III: L2–L4 outcomes by model and condition. L2 is trivially adopted by all models (three at $n = 20$; `gpt-4o-mini` pilot $n = 5$); L3 reveals a model-capability threshold (`gpt-4o-mini` fails, others succeed); L4 (example-only) defeats all models despite L3 success. Rows with $n = 5$ are pilot results; Wilson 95% CIs shown.

Model	Level	Delivery	Pass	Failure mode
`gpt-4o-mini`	L2	Expl. grammar	5/5†	None (adopted)
`gpt-4o`	L2	Expl. grammar	20/20	None (adopted)
`gpt-4.1-mini`	L2	Expl. grammar	20/20	None (adopted)
`gpt-4.1-nano`	L2	Expl. grammar	20/20	None (adopted)
`gpt-4o-mini`	L3	Expl. rule	0/5†	Prior dominance
`gpt-4o`	L3	Expl. rule	20/20	— (success)
`gpt-4.1-mini`	L3	Expl. rule	20/20	— (success)
`gpt-4.1-nano`	L3	Expl. rule	20/20	— (success)
`gpt-4o`	L4	Examples only	0/5†	Pattern blindness
`gpt-4.1-nano`	L4	Examples only	0/50	Pattern blindness

†Pilot ($n = 5$); Wilson 95% CI: [57%,100%] for 5/5, [0%,43%] for 0/5.

each (SIR=0.0; Wilson 95% CI: [84%,100%] for adoption rate); `gpt-4o-mini` adopted L2 in a pilot evaluation ($n = 5$, 5/5; Wilson CI: [57%,100%]). Across all four evaluated models, SIR=0.0, confirming that structural novelty alone does not create LLM resistance.

C. L3: Capability threshold at semantic rule-following

With an explicit semantic inversion rule, `gpt-4o-mini` fails completely (0/5, PPR=1.0): it generates standard Python despite instruction. By contrast, `gpt-4o`, `gpt-4.1-mini`, and `gpt-4.1-nano` all succeed (20/20, PPR=0.0). The gap between `gpt-4o-mini` (0/5) and `gpt-4o` (20/20) is significant (Fisher's exact $p < 0.0001$), revealing a model capacity threshold for following semantic meta-rules.

Key contrast: L3 success $\not\Rightarrow$ L4 success. `gpt-4.1-nano` achieves 20/20 at L3 but 0/50 (PPR=1.0) at L4—the same model that correctly follows an explicit inversion rule is completely blind to the same rule embedded only in examples. This within-model contrast directly isolates pattern blindness as a failure mode independent of model capacity to follow rules when stated explicitly.

D. L4: Pattern blindness across all tested models

Without explicit rule, `gpt-4o`—which succeeds at L3—fails at L4 (PPR=1.0), producing canonical Python and ignoring example-embedded inversion cues. Figure 4 shows the contrast.

L4 ablation (n=50 per model). We scaled the L4 evaluation across seven models, all at $n = 50$ (`gpt-4o`, `gpt-4o-mini`, `gpt-4.1`, `gpt-4.1-mini`, `gpt-4.1-nano`, `gpt-5.4-mini`, `o4-mini`). Results are shown in Table IV. All seven models achieve 0% pass rate, confirming that pattern blindness is robust across model families and including the reasoning model `o4-mini`. (`gpt-5.4-mini`: 0/42 valid runs, $n_{\text{eff}} = 42$; see footnote.) PPR values range from 0.38 (`gpt-4.1-mini`) to 1.0 (`gpt-4.1`, `gpt-4.1-nano`, `gpt-5.4-mini`); `o4-mini` shows PPR=0.92. (`gpt-4.1-mini`'s 0.38 overall PPR reflects variant heterogeneity, not a uniformly lower prior pull; variant-level PPR spans 0.00–1.00, with certain variants achieving full inversion while others revert completely—see Table IV for per-variant values.) Notably, `gpt-4.1` and `gpt-4.1-nano` show PPR=1.0 across *all* 50 runs and all variants, indicating complete prior dominance with no evidence of inversion detection. The `o4-mini` result revises an earlier preliminary finding (1/20 at $n = 20$): at the full scale of $n = 50$, `o4-mini` also achieves 0/50, indicating that chain-of-thought reasoning does not provide a reliable pathway to overcoming pattern blindness at this scale.

TABLE IV: L4 ablation (PPR = Python Prior Rate). Seven core models: $n = 50$ each (5 variants $\times$ 10 runs); `gpt-5.4-mini` $n_{\text{eff}} = 42$ (8 HTTP errors in Var. E, excluded). Pass rate = 0 for all models. `o4-mini` revises $n = 20$ preliminary (1/20). §E24 frontier extension: `gpt-5.4` and `o3` also tested at Var. A, $n = 10$: both 0/10 pass, PPR=1.00.

Var.	Description	gpt-4o	gpt-4o-mini	gpt-4.1	gpt-4.1-mini	gpt-4.1-nano	gpt-5.4-mini	o4-mini
A	2ex-baseline	0.40	**0.10**	1.00	**0.10**	1.00	1.00	0.80
B	3ex-extra	1.00	1.00	1.00	**0.00**	1.00	1.00	1.00
C	2ex-no-comment	0.90	1.00	1.00	0.80	1.00	1.00	0.90
D	2ex-swapped	0.70	1.00	1.00	**0.00**	1.00	1.00	1.00
E	3ex-diff-domain	0.50	1.00	1.00	1.00	1.00	N/A†	0.90
Overall PPR		0.70	0.82	1.00	**0.38**	1.00	1.00‡	0.92
Pass rate		0/50	0/50	0/50	0/50	0/50	0/42†	0/50

†`gpt-5.4-mini` var-E: 8 HTTP errors; $n_{\text{eff}} = 2$; excluded from per-variant PPR.
‡Overall PPR for `gpt-5.4-mini` computed over 4 valid variants (A–D); var-E excluded.
§E24: `gpt-5.4` 0/10 (PPR=1.00), `o3` 0/10 (PPR=1.00) on Var. A only; see text.

E24: Frontier models on L4 Fibonacci (strict example-only, $n = 10$). To close the gap identified in §V-N—where `gpt-5.4` and `o3` were only tested on hard tasks—we ran both frontier models on the identical Variant A condition used in Table IV (strict example-only, Fibonacci, $n = 10$). Result: `gpt-5.4` 0/10 (PPR=1.00) and `o3` 0/10 (PPR=1.00). Both frontier models show complete prior dominance on L4 Fibonacci, confirming that pattern blindness persists even in the latest dense and reasoning models when the task has a deeply entrenched Python prior. This result directly confirms the core claim: *L4 pattern blindness on prior-entrenched tasks is not defeated by frontier models as of 2026-03-25*.

Variant E domain-shift anomaly (`gpt-4.1-mini`). The most striking per-variant result is `gpt-4.1-mini`'s complete inversion across variant E: PPR=0.00 on variants B (3ex-extra) and D (2ex-swapped)—meaning it never generates the Python prior—yet PPR=1.00 on variant E (3ex-diff-domain, where examples use a *different* semantic domain). The structural difference is that variants B and D use the same Fibonacci-adjacent domain for all examples, while variant E uses cross-domain examples (e.g., sorting or arithmetic comparisons). We hypothesize that *domain shift in examples weakens the model's pattern extraction signal*: when examples share the same function context as the target task, the model's in-context pattern extractor can align examples to the task template; when examples are drawn from a different domain, the learned pattern loses its grounding in the target task structure, causing the model to revert to the Python prior. This remains a hypothesis; controlled cross-domain vs. in-domain replication at larger n would be needed to confirm it. Notably, variant A (2ex-baseline, in-domain) achieves PPR=0.10 while variant B (3ex,

✓ Panel A: L3 — explicit rule (`gpt-4o`) [CORRECT]

System: "if-blocks run when condition is FALSE."

`gpt-4o` output **(CORRECT):**

```
if n > 1:          # inverted
    return n
return fib(n-1) + fib(n-2)
```

× Panel B: L4 — examples only (`gpt-4o`) [INCORRECT]

User: [2 worked examples with inverted semantics, no explicit rule]

`gpt-4o` output **(INCORRECT):**

```
if n <= 1:         # standard Python
    return n
return fib(n-1) + fib(n-2)
```

Fig. 4: Prior dominance vs. pattern blindness isolated on the same model (`gpt-4o`). **A**: `gpt-4o` correctly inverts logic with explicit rule (L3). **B**: The same model reverts to Python prior when the rule is embedded only in examples (L4). This contrast isolates pattern blindness as a failure mode distinct from prior dominance.

in-domain) achieves PPR=0.00: example count and domain *jointly* modulate extraction quality in this model—additional in-domain examples may partially suppress the Python prior without enabling full inversion, while domain shift reinvigorates it.

E. L4 Multi-task Generalization

Protocol note: This section uses *partially-annotated* prompts (inline comments partially state the inversion rule), which differ from the strict *example-only* prompts in §V-D. Results in Table V are **not directly comparable** to Table IV; comparisons across these sections should account for this prompt protocol difference.

To assess whether pattern blindness extends beyond Fibonacci, we evaluated all seven models on three tasks: T1 (`fibonacci`), T2 (`is_sorted`), and T3 (`count_vowels`). The multi-task prompt uses inline code comments that partially annotate the inversion rule (*e.g.*, `# runs when a > b is FALSE, i.e., a <= b`); this is a *partially-annotated L4* condition—less explicit than L3 (no stated rule), but more informative than the pure example-only variants used in the ablation. `gpt-4o`, `gpt-4o-mini`, `gpt-5.4-mini`, and `o4-mini` were evaluated at $n = 10$ or $n = 15$ per task; `gpt-4.1-mini` at $n = 20$ per task; `gpt-4.1` and `gpt-4.1-nano` at $n = 10$ per task (255 initial runs). We further scaled `gpt-4o` T3 and `gpt-4.1` T2 to $n = 50$ each to resolve preliminary ambiguities, and extended `o4-mini` T2/T3 to $n = 20$ each, yielding 345 total multi-task runs. Results are in Table V.

TABLE V: L4 multi-task summary (**345 runs total**). † Partially-annotated prompts; **NOT directly comparable to Table IV** (strict example-only). Green = pass, red = fail. Dominant failure mode per row in rightmost column. Full PPR values in Appendix X-D.

Model	T1: fib	T2: sorted	T3: vowels	Dominant failure
`gpt-4o`	0/10	**10/10**	0/50[a]	T1: prior dom.; T3: op. sub. (94%)
`gpt-4o-mini`	0/10[b]	0/10	0/10	template fragmentation (T1); prior dom.
`gpt-4.1`	0/10	1/50[c]	0/10	prior dominance
`gpt-4.1-mini`	0/20	**20/20**	0/20	T1: prior dom.; T2: extractable
`gpt-4.1-nano`	0/10	0/10	0/10	prior dominance
`gpt-5.4-mini`	0/10	0/10	0/10	prior dominance (all tasks)
`o4-mini`	**8/15**[d]	0/20	0/20	T1: partial extraction (reasoning); T2/T3: prior dom.

[a] OpSub 47/50 (94%, CI: [83%,99%]). [b] Template fragmentation.
[c] $n = 50$ scale-up; Wilson CI: pass [0%,10.5%]; Fisher's $p<0.0001$ vs. `gpt-4.1-mini`.
[d] Token budget corrected (400→5000 tokens); see §V-D for details.

Seven findings emerge.

(1) **T2 is extractable only by mid-to-large models**: `gpt-4o` (10/10) and `gpt-4.1-mini` (20/20) pass T2 fully, while `gpt-4.1` (1/50, PPR=0.42), `gpt-4.1-nano` (0/10), `gpt-5.4-mini` (0/10), and `gpt-4o-mini` (0/10) fail. This reveals a non-monotonic capacity threshold: T2 extractability does not scale linearly within or across model families.

(2) **T1 (Fibonacci) is resistant for all non-reasoning models**: six of seven models fail T1, with high or full PPR (PPR=1.0 for `gpt-4o`, `gpt-4.1`, `gpt-4.1-mini`, `gpt-5.4-mini`; PPR=0.2 for `gpt-4.1-nano`; template fragmentation for `gpt-4o-mini`). The reasoning model `o4-mini` provides a partial exception under the partially-annotated prompt (see Finding 5 for detail).

(3) **Operational substitution dominates T3**: the scale-up (`gpt-4o` T3, $n = 50$) reveals operational substitution in 47/50 runs (94%; Wilson 95% CI: [83%, 99%])—the model applies arithmetic transformations (`count -= 1; return -count`) to achieve correct output without adopting the L4 semantic rule. The preliminary $n = 10$ result (1 pass) was a statistical artifact; at $n = 50$, `gpt-4o` achieves 0/50 pass (Wilson 95% CI: [0%, 7.1%]). `gpt-5.4-mini` T3 also exhibits operational substitution (PPR=0.0, all 10 runs), confirming this as a structurally distinct, frequent failure mode.

(4) **Model capacity determines T2 extractability**: `gpt-4o-mini`, `gpt-4.1`, `gpt-4.1-nano`, and `gpt-5.4-mini` all fail T2, while `gpt-4o` and `gpt-4.1-mini` succeed—consistent with the capacity threshold observed at L3.

(5) **Reasoning partially mitigates pattern blindness under partial annotation**: `o4-mini` achieves 8/15 on T1 under the partially-annotated prompt (after token budget correction), but fails T2 and T3 (0/20 each at extended $n = 20$; overall 8/55 across T1+T2+T3). Under the strict example-only ablation (§ V-D), `o4-mini` achieves 0/50—chain-of-thought reasoning alone cannot overcome pure pattern blindness, but reasoning models can exploit partial annotation cues that dense models miss.

Methodological note on `o4-mini` (token budget correction). An initial `o4-mini` T1 run used

max_completion_tokens=400, which truncated the model's chain-of-thought before it could complete the inversion task; this yielded 0/15 pass (consistent with pattern blindness). Upon correcting the budget to max_completion_tokens=5000, o4-mini achieved 8/15 pass—representing a finding revision. This correction reveals that reasoning models can exploit partial annotation cues inaccessible to dense models, but only when given sufficient token budget to complete the reasoning chain. The original 0/50 result in the L4 ablation (strict example-only, corrected budget) confirms that this effect is annotation-contingent: token budget alone does not overcome pure pattern blindness.

(6) **Within-family inconsistency is confirmed**: gpt-4.1 T2 scale-up ($n = 50$) confirms the capacity gap is systematic—1/50 pass (Wilson 95% CI: [0%, 10.5%]) vs. gpt-4.1-mini 20/20 (Fisher's exact $p < 0.0001$). Fine-tuning or RLHF choices appear to play a larger role than raw scale in determining L4 pattern extraction capability.

(7) **gpt-5.4-mini shows uniform prior dominance across all tasks**: 0/30 with T1 PPR=1.0, T2 PPR=1.0, and T3 operational substitution pattern, consistent with its L4 ablation PPR=1.0 on Fibonacci variants.

F. Three Failure Modes

- **Prior Dominance** (L1, L3-weak, L4-T2-mini, L4-T1-all): stated or implicit rule is overridden by pretrained token distribution.
- **Pattern Blindness** (L4-T1): model cannot extract the semantic rule from examples alone, even when the same model follows the rule when stated explicitly (L3). T1 (Fibonacci) shows universal pattern blindness across all models.
- **Operational Substitution** (L4-T3-gpt-4o, $n = 50$): model compensates via arithmetic output transformation rather than adopting the semantic inversion rule, indicating goal awareness without rule internalization. Confirmed as dominant at 94% of T3 gpt-4o runs (Wilson 95% CI: [83%, 99%]).

We also observe *template fragmentation* in gpt-4o-mini T1: the model partially disrupts its standard template but does not achieve semantic inversion, yielding an incorrect but structurally novel output. We treat this as a sub-variant of prior dominance rather than a distinct fourth mode, given the small sample ($n = 10$); further study at larger n may warrant reclassification.

Task difficulty interacts with model capacity to determine which failure mode dominates. T2 (is_sorted) is extractable by mid-to-large models; T1 (Fibonacci) is universally resistant; T3 (count_vowels) triggers operational substitution as the dominant escape strategy for capable models.

Failure taxonomy across 419 failed L4 runs (ablation + multitask T1; 427 total pool including 8 o4-mini T1 passes): Type-I (exact Python prior) accounts for ≈83.1% of failures, Type-II (partial attempt) for ≈12.6%, and Type-III (operational substitution) for ≈4.3%. Type-I dominates in gpt-4.1 and gpt-4.1-nano (100% of failures); Type-II concentrates in gpt-4.1-mini (44%) and gpt-4o-mini (32%); Type-III appears primarily in gpt-4o (25% of its failures).

Population note: The 427 pooled runs consist of 342 valid strict example-only ablation runs (all failures; gpt-5.4-mini $n_{\text{eff}} = 42$ due to 8 HTTP errors, counted at 342 not 350) and 85 multitask T1 runs (8 passes by o4-mini, 77 failures). Total failures: 419. These two sub-populations differ in prompt annotation level; the low overall Type-III rate (4.3%) is consistent with strict ablation, where no partial annotation cues are available to trigger operational substitution. In the multi-task sub-population, Type-III is observed primarily in gpt-4o T3 at $n = 50$ (94% operational substitution rate), but that data is counted in T3 (not T1) and is therefore excluded from the 427 pooled taxonomy. A full split analysis per prompt protocol is deferred to future work.

G. CoT Ablation: Reasoning About Inversion Is Not Sufficient

We tested whether chain-of-thought (CoT) prompting overcomes pattern blindness by adding "think step by step about the pattern first" to the L4 Variant A prompt ($n = 20$ per condition, four models: gpt-4o, gpt-4.1-mini, gpt-4.1-nano, o4-mini).

Result: CoT does not reliably improve pass rate. Three of four models achieve 0/20 in both CoT and no-CoT conditions (Wilson 95% CI: [0%,16%]); o4-mini achieves 1/20 under CoT (5%; Wilson CI: [0.9%,24%]; Fisher's exact vs. 0/20 no-CoT: $p = 1.0$, not significant). However, CoT consistently changes *reasoning behavior* across all models:

- gpt-4o: CoT reduces PPR from 1.00 to 0.15 and raises *mentions_inversion* to 100%—the model explicitly *describes* the inverted pattern but still generates standard Python code.
- gpt-4.1-mini: CoT reduces PPR from 0.40 to 0.00 with 100% mentions_inversion rate—same reasoning–generation dissociation.
- gpt-4.1-nano: CoT reduces PPR from 1.00 to 0.20 and raises *mentions_inversion* to 100%—the nano model shows the strongest PPR reduction under CoT, yet all 20 outputs still fail to implement the inversion correctly.
- o4-mini: CoT reduces PPR from 0.95 to 0.65 and raises *mentions_inversion* to 35%. This is the most notable reasoning-model CoT result: while dense models reach near-perfect *mentions_inversion* (100%) yet still fail, o4-mini reaches only 35% mentions at $n = 20$, suggesting a dissociation between CoT-prompting and chain-of-thought use in the reasoning model. The no-CoT baseline (PPR=0.95 at Variant A, $n = 20$) is slightly higher than the overall L4 ablation PPR=0.92 for o4-mini (which averages all five variants); this is expected since different variants show different PPR values. The CoT comparison should be interpreted within Variant A only.

This reveals a *reasoning–generation dissociation* that generalizes across four model families: models can articulate the semantic inversion rule when prompted to reason, yet their code generation defaults to the Python prior regardless. Pattern blindness is thus not a failure to *detect* the pattern but a

failure to *override* the prior during generation. The finding is robust: 160 total CoT runs (40 per model, 2 conditions) across dense and reasoning architectures yield zero reliable pass rate improvement.

CoT vs. partial annotation for `o4-mini`. A notable dissociation emerges when comparing the CoT ablation result to the multi-task result for `o4-mini`: CoT instruction on strict Variant A yields 1/20 (5%), while partial annotation (inline comment hints) without CoT instruction yields 8/15 (53%, Table V). This suggests that for reasoning models, partial rule disclosure is more effective than meta-cognitive prompting alone: CoT helps dense models articulate the inversion pattern but does not supply the rule content; partial annotation provides actual rule content that reasoning models can exploit. The implication for resistance design is that partial annotations—not just implicit example delivery—are the key vulnerability for reasoning models.

H. Open-Weight Models: Preliminary L1 and L4 Results

To assess whether the failure modes generalize beyond the OpenAI model family, we conducted preliminary evaluations on two open-weight models via Groq API: `Llama-3.3-70B-Versatile` and `Qwen3-32B`.

L1 (ctx-pack, $n = 20$): `Llama-3.3-70B` achieves KLR=0.23, nearly matching the OpenAI best (`gpt-4.1-mini` ctx-pack, 0.21). `Qwen3-32B` achieves KLR=1.00: all 20 outputs contained at least one native Python keyword with no valid code skeleton—the worst L1 result across all tested models. The inter-model gap (0.23 vs. 1.00) exceeds the intra-family spread of `gpt-4.x` models (0.21–0.69), indicating that instruction-following fine-tuning rather than parameter count governs alias compliance.

L4 Variant A ($n = 20$): Both open-weight models fail. `Llama-3.3-70B`: 0/20 pass, PPR=0.05 (1/20 strict Python prior; remaining 19 outputs applied standard control flow without inversion). `Qwen3-32B`: 0/18 pass ($n_{\text{eff}} = 18$; 2 HTTP errors), PPR=1.00 (complete prior dominance). Neither model triggered inversion detection (noticed_inversion=0). These results extend the pattern blindness finding beyond OpenAI models.

Limitations and extended runs: Initial pilot ($n = 20$, Variant A only) confirmed pattern blindness for both open-weight models. Full 5-variant ablation ($n = 50$ attempted) was subsequently run for both models. `Qwen3-32B` completed 50/50 valid runs (0 HTTP errors): 0/50 pass, PPR=1.00 across all variants—consistent with pilot results, confirming complete prior dominance. `Llama-3.3-70B` achieved 25/50 valid runs (25 HTTP errors due to Groq rate limiting; all five variants represented, $n_{\text{eff}} = 5$ per variant): 0/25 pass. Per-variant PPR values vary substantially: Var. A=0.00, Var. B=0.00, Var. C=0.60, Var. D=1.00, Var. E=1.00 (overall PPR=0.52). At $n_{\text{eff}} = 5$ per variant, Wilson 95% CIs are wide: PPR=0.00 has CI [0%, 43%]; PPR=0.60 has CI [17%, 93%]; PPR=1.00 has CI [57%, 100%]. These intervals overlap substantially—statistical claims about variant-level differences are not warranted at this sample size. This per-variant *pattern* suggests a qualitatively different failure mode from both prior dominance (PPR≈1.0) and Qwen's complete prior dominance: `Llama-3.3-70B` may apply standard control flow without semantic inversion (low PPR on Variants A/B) but revert to strict Python prior on Variants D/E (domain-shifted examples). This potentially represents a third open-weight failure mode—*context-sensitive prior reversion*—but the per-variant pattern cannot be statistically confirmed at $n_{\text{eff}} = 5$; full replication is deferred to future work.

TABLE VI: E28: Annotation density gradient for `o4-mini` on T1 Fibonacci ($n = 10$ per level, 40 total runs; $n_{\text{eff}} = 9$ at 100% due to 1 HTTP error). PPR = prior-pull rate (fraction of `fail_prior` among all non-pass runs). Pass rate and PPR are both reported.

Density	Annotated/Total	Pass	Fail-Prior	Fail-Other	Pass Rate	PPR
25%	2/8	3	1	6	0.30	0.10
50%	4/8	6	1	3	**0.60**	0.10
75%	6/8	2	5	3	0.20	0.50
100%†	8/8	1	2	6	0.11	0.22

† $n_{\text{eff}} = 9$ (1 HTTP error); pass rate = 1/9.

E26/E27: L3 diagnostic for `Qwen3-32B` ($n = 5$ pilot, $n = 20$ replication). To disambiguate whether `Qwen3-32B`'s L4 failure (PPR=1.00) reflects (a) pattern extraction failure (*L4-specific*, same as OpenAI models) or (b) general instruction-following failure for semantic meta-rules, we ran $n = 5$ L3 evaluations (E26) and subsequently $n = 20$ full replication (E27) on `Qwen3-32B` (explicit rule: "if-blocks execute when the condition is FALSE," Fibonacci task). E26 result ($n = 5$): 0/5 pass, PPR=0.80 (1/5 showed a structurally distinct non-pass failure). E27 result ($n = 20$, 2026-03-25): 0/20 pass, PPR=0.40 (8/20 exact Python prior; 12/20 FAIL-other). No valid passes at $n_{\text{eff}} = 20$; Wilson 95% CI: [0%, 17%]. The PPR=0.40 is *lower* than the E26 pilot (0.80), reflecting that a substantial fraction of failures are structurally distinct from the pure Python prior—the model often produces internally inconsistent code that attempts partial rule application but fails, rather than simply returning the standard Fibonacci template. This is consistent with *confused partial instruction-following*, not clean prior dominance. Unlike `gpt-4.1` (E25: 10/10 on L3 T2), `Qwen3-32B` fails to apply the semantic inversion rule even when stated explicitly across $n = 20$ trials—confirming *general semantic meta-rule failure* as a systematic property, not a small-n artifact. This represents a qualitatively different failure mode from all tested OpenAI models: `Qwen3-32B` cannot reliably follow semantic inversions via prompt instructions alone, regardless of delivery (explicit vs. implicit).

I. E28: Annotation Density Gradient for `o4-mini` (Pilot)

This section reports the single-model pilot; the full three-model replication (E34) is in §V-M. The E28 experiment reveals that annotation density has a **non-monotonic** effect on `o4-mini`'s ability to extract the L4 inversion pattern. Results are in Table VI.

Key finding—optimal density at 50%: Pass rate peaks at 50% annotation density (6/10, PPR=0.10). Below this, at

25% density, pass rate drops to 3/10—insufficient annotation signal for the reasoning model's chain-of-thought. Above 50%, at 75% annotation, pass rate collapses to 2/10 and PPR *rises* sharply to 0.50—a paradoxical result where more annotation triggers stronger prior conflict. At 100% density, pass rate is lowest (1/9 valid runs; 1 HTTP error), though PPR partially recovers to 0.22.

Interpretation—annotation saturation induces prior conflict: The non-monotonic curve is consistent with a two-competing-signal model. At low density (25%), the reasoning model has insufficient annotation cues to construct the inversion rule: it either succeeds via independent pattern recognition (30%) or fails via non-prior confusion (60%). At 50% density, annotation signal is sufficient for reasoning chain completion without overwhelming the Fibonacci prior. At 75%+, the dense annotation makes the inversion semantics *salient* to the model's representation, which simultaneously activates the conflicting Python prior—triggering more frequent prior-dominant failures (PPR=0.50). This suggests a retrieval-interference mechanism: annotation cues strong enough to activate the task schema also activate the competing prior schema.

Design implication: For prior-entrenched tasks, resistance via annotation density is robust in a narrow band. Over-annotated prompts are *not* strictly harder; they may in fact partially self-defeat by activating the prior they target. Optimal confusion language design should calibrate annotation density to remain below the prior-activation threshold—approximately 50% of examples for reasoning models on Fibonacci-class tasks.

Connection to CoT–annotation dissociation: The E28 results reinforce Finding 5's CoT–annotation gap. The prior multi-task result (8/15 pass at partial annotation) used a fixed partial-annotation regime that approximates the 50% optimal; the CoT ablation (1/20 pass) added meta-cognitive instruction without annotation content. E28 maps the mechanism: annotation works by supplying inversion-rule content at a density that clears the reasoning threshold without triggering prior conflict. CoT alone cannot cross this boundary because it adds reasoning scaffolding but no new rule content.

J. E31: Prior Entrenchment Score via Logprobs

To directly measure prior strength, we queried the logprobs API at the key decision point of each task's base case (the token immediately after `return`). For each model, we recorded the top-5 token probabilities and the rank/probability of the canonical Python token. Table VII summarizes the results.

Key finding—PES alone does not predict L4 resistance. Both `fib` and `merge_sort` show near-identical PES ($\geq$0.9996), yet their L4 outcomes diverge dramatically: fib is universally resistant (0/350 pass in ablation), while merge_sort is solvable by all tested models in the hard-task condition (Table XI). This dissociation rules out a simple "prior strength $\rightarrow$ resistance" mapping and supports the hypothesis that I/O specification transparency—not prior magnitude alone—is the critical moderating variable (§V-N). The PES metric quantifies the *strength* of the prior that must be overridden, but the model's ability to override it depends on whether the task's I/O examples make the target behavior transparent.

TABLE VII: E31: Prior Entrenchment Score (PES) — probability of the canonical Python token at the base-case decision point. Models: `gpt-4o`, `gpt-4.1-mini`. `gpt-4o-mini` results excluded (model emits markdown formatting tokens rather than code tokens at this position). PES = P(canonical token | prompt prefix).

Task	Canonical token	gpt-4o	gpt-4.1-mini
fib (T1)	`n`	0.9996	1.0000
merge_sort (H1)	`arr`	0.9999	1.0000
is_sorted (T2)	`True`	0.9736	0.9985
count_vowels (T3)	`s`	0.1606	1.0000

TABLE VIII: E32: Cross-task transfer — pass rate on L4 Fibonacci after exposure to inverted code from other tasks. Baseline: strict L4 example-only (Table IV) yields 0/50 for all models. Wilson 95% CIs shown. `gpt-4o` scaled to $n = 30$ (E32a); others at $n = 10$. E32-ctrl: control condition using normal (non-inverted) sort examples, `gpt-4o`, $n = 10$.

Model	Pass	Pass rate	Wilson 95% CI	Note
`gpt-4.1-mini`	10/10	**100%**	[72%,100%]	Perfect
`gpt-4o`	16/30	53%	[36%,70%]	Partial t
`o4-mini`	1/10	10%	[2%,44%]	Minima
E32-ctrl (`gpt-4o`, normal examples)	0/10	0%	[0%,31%]	No tran

`is_sorted` (T2) shows slightly lower PES (0.9736 for `gpt-4o`), which may partly explain its extractability by capable models—but PES=0.9736 is still very high, suggesting that the $\sim$0.03 PES gap is not the primary driver of the T1–T2 pass rate difference (0% vs. 100% for `gpt-4o`).

A secondary observation: `gpt-4o`'s low T3 PES (0.1606, `count_vowels`) contrasts with its high operational substitution rate on T3 ($n = 50$: 94% OpSub). This suggests that a weaker prior (lower PES) does not simply yield L4 compliance; instead, the model explores alternative generation paths that still avoid semantic inversion, manifesting as arithmetic workarounds rather than canonical Python. The PES–operational-substitution relationship warrants future mechanistic study.

K. E32: Cross-Task Transfer Breaks L4 Resistance

A critical question for assessment design: can a model crack L4 Fibonacci if it first sees inverted semantics demonstrated on *other* tasks? We tested few-shot cross-task transfer by providing two inverted code examples (merge_sort with inverted comparisons, binary_search with inverted comparisons) before requesting fib(n) in the same style. `gpt-4.1-mini` and `o4-mini` were evaluated at $n = 10$; `gpt-4o` was scaled up to $n = 30$ (two scale-up rounds) to tighten the confidence interval on its partial-transfer result. Temperature=0.7 for all. A **control condition** (E32-ctrl: ascending sort with *normal*, non-inverted examples preceding fib, `gpt-4o`, $n = 10$) was run to confirm that the transfer effect is specific to inversion semantics, not merely cross-task code exposure.

TABLE IX: E35 (Redesigned I/O transparency): pass rate on IO-A (strong Fibonacci prior, transparent mapping) vs. IO-B (weak prior, opaque I/O only — no algorithm description). $n = 10$ per cell, temperature=0.7.

Model	IO-A Pass	IO-A Rate	IO-B Pass	IO-B Rate
gpt-4o	10/10	**100%**	0/10	**0%**
gpt-4.1-mini	10/10	**100%**	0/10	**0%**
gpt-4o-mini	5/10	50%	0/10	**0%**

Cross-task contamination defeats L4. The same models that achieve 0/50 under strict L4 example-only delivery (Table IV) achieve 53–100% pass when primed with inverted code from unrelated tasks. gpt-4.1-mini transfers perfectly (10/10; Wilson 95% CI: [72%,100%]), consistent with its strong pattern extraction on T2 (20/20 in multi-task). gpt-4o achieves 16/30 at scale-up (Wilson 95% CI: [36%,70%])—partial transfer confirmed; Fisher's exact vs. gpt-4.1-mini: $p = 0.008$ ($p < 0.01$), indicating a significant gap in transfer fidelity. o4-mini—which relies on reasoning chains rather than pattern mimicry—achieves only 1/10 (Wilson 95% CI: [2%,44%]), suggesting that systematic reasoning about inverted semantics does not facilitate transfer in the way that code-example mimicry does. The control condition (E32-ctrl: 0/10) confirms that the transfer effect is specific to *inverted* semantics: exposure to non-inverted cross-task code leaves L4 resistance fully intact, ruling out mere cross-task priming as an explanation.

Assessment design implication. L4 resistance depends on *isolation*: if students include inverted code from other tasks in the same prompt (e.g., copying solved examples from a previous assignment), the model can transfer the inversion pattern cross-task. Instructors should ensure that L4 assignments do not provide cross-task inversion examples within the same prompt context. This finding also suggests that L4 resistance is a property of *prompt-level pattern isolation*, not just task-level prior entrenchment.

L. E33: I/O Transparency — Prior vs. Opaque Mapping

We conducted two experiments to isolate I/O transparency from prior strength. The original E33 design was confounded: the IO-B (opaque_counter) prompt disclosed explicit algorithmic steps (XOR, shift, conditional add/subtract), making both tasks procedurally transparent; all three models (gpt-4o, gpt-4.1-mini, gpt-4o-mini) achieved 10/10 on both tasks. A replication (E33v2) with a different model set (gpt-4o, gpt-4.1-mini, o4-mini) confirmed this result: all models again achieved 10/10 on both IO-A and IO-B, indicating the confound is robust across model generations when algorithmic steps are disclosed. To properly test I/O opacity as an independent variable, we ran E35 (**redesigned**): IO-B prompts were replaced with pure input/output pairs ($n = 8$ examples), with no algorithm disclosed. IO-A remained identical to E33 (explicit "reverse Fibonacci" semantics, Fibonacci prior).

TABLE X: E34: Annotation density curve — pass rate at each density level ($n = 10$ per cell, 150 total runs). MAT = Minimum Annotation Threshold (lowest level where pass rate >50%). The 0% condition includes an inverted *code example* (not I/O-only); it is therefore not comparable to the strict L4 ablation.

Model	0%†	25%	50%	75%	100%	MAT
gpt-4o	5/10	0/10	1/10	2/10	3/10	None
gpt-4.1-mini	**8/10**	1/10	5/10	4/10	1/10	0%†
o4-mini	**7/10**	3/10	6/10	6/10	1/10	0%†

†Code-example (mimicry) baseline; qualitatively distinct from verbal annotation (25–100%).

Opaque I/O creates universal resistance. When models receive only input/output pairs with no algorithmic description, all three models fail completely on IO-B (0/30 pooled; Wilson 95% CI: [0%,12%]). By contrast, IO-A—which provides explicit naming cues ("reverse Fibonacci") that activate a well-known prior—achieves 100% for both larger models. The IO-A/IO-B contrast (100% vs. 0% for gpt-4o/gpt-4.1-mini) demonstrates that I/O opacity alone is sufficient to induce resistance, independent of prior strength: the opaque_counter function has no deep training prior, yet models cannot infer or generalize the underlying bijection from 8 examples to held-out inputs. The gpt-4o-mini IO-A result (5/10) reveals an additional confound: even a transparent I/O mapping fails 50% when the *Fibonacci prior* interferes with the reversed sequence ordering, demonstrating bidirectional prior effects.

Design implication. L4 resistance is not purely prior-strength-driven; it can also be induced by I/O-opaque task designs for which no training distribution provides a shortcut. This suggests a complementary pathway to NTP-resistant assignment design: rather than requiring the model to violate a known prior (L4), instructors can use opaque I/O specifications that prevent prior activation entirely.

M. E34: Annotation Density Curve (Multi-Model)

E28 established a non-monotonic annotation density curve for o4-mini alone. E34 extends this to three models across five density levels (0%, 25%, 50%, 75%, 100%) on L4 Fibonacci with inverted comparisons ($n = 10$ per level per model, 150 total runs). **Important:** the 0% level is a *code-example (mimicry) baseline*—it includes an inverted code example rather than I/O-only delivery as in the strict L4 ablation, making it qualitatively distinct from the verbal annotation range (25–100%); 100% provides the full explicit rule (L3-equivalent). Results are in Table X.

Key finding—pattern mimicry via code examples outperforms verbal annotation at all density levels. All three models show highest pass rates at 0% annotation density (which includes an inverted *code* example, a qualitatively distinct condition from verbal annotation at 25–100%), *not* at 100% (full explicit rule). This is the opposite of the naive expectation that more explicit instructions improve compliance. Within the verbal annotation range (25–100%), pass rates are non-monotonic: (1) sharp collapse at 25% (single comment hint disrupts code-example mimicry without providing sufficient

TABLE XI: L4 hard-task results ($n = 10$ per cell, 120 total runs). All tasks use example-only delivery with 2 worked I/O examples. Pass = output matches inverted specification. Models include `gpt-5.4` and `o3` (not tested in main L4 ablation). Wilson 95% CIs at $n = 10$: 10/10 → [72%,100%]; 0/10 → [0%,28%]; `gpt-4o` H1 (4/10) → [17%,69%]—interpret with caution.

Model	H1: merge_sort	H2: bin_search	H3: bfs	Overall
`gpt-4o`	4/10	10/10	10/10	80%
`gpt-5.4`	10/10	10/10	10/10	**100%**
`o3`	10/10	10/10	10/10	**100%**
`o4-mini`	10/10	10/10	10/10	**100%**

rule content), (2) partial recovery at 50–75%, and (3) collapse at 100% (full rule triggers prior conflict). The overall pattern distinguishes two qualitatively different mechanisms: code-example mimicry (effective at 0%) and verbal rule induction (weak and non-monotonic across 25–100%).

Pattern mimicry > rule following. The 0% condition succeeds because the inverted code example provides a direct structural template for mimicry—the model copies the comparison operator pattern without needing to understand the underlying rule. Adding explicit annotations paradoxically *hurts* because it shifts the model from pattern mimicry mode to rule-following mode, where the Python prior creates stronger interference. This is consistent with E32's cross-task transfer finding: code examples are more effective than verbal rules for inducing inverted behavior.

Contrast with strict L4 ablation. The 0% pass rates here (50–80%) contrast with the strict L4 ablation's 0/350. The critical difference: E34's 0% prompt shows an inverted *code* example (`if n >= 0`), while the strict L4 ablation shows only I/O examples and worked solutions without explicitly inverted comparison operators in the prompt code. This confirms that L4 resistance depends on *what* examples show, not just whether a rule is stated.

Convergence with E28 (`o4-mini`). E34's `o4-mini` results are directionally consistent with E28: both experiments agree at 50% annotation density (E28: 6/10; E34: 6/10). The apparent "peak at 0%" in E34 is due to the qualitatively distinct code-example baseline—E28 did not include a 0% code-example condition, so the two experiments are directly comparable only in the 25–100% verbal annotation range, where E34 shows the same non-monotonic pattern (50%/75% local peak, collapse at 100%) consistent with E28. Both experiments confirm that annotation saturation hurts reasoning models on prior-entrenched tasks.

N. L4 Hard-Task Extension: Task Complexity vs. Prior Entrenchment

To test whether more complex algorithms amplify or defeat pattern blindness, we evaluated three hard tasks (H1: descending merge sort, H2: inverted binary search, H3: reversed BFS) across four models including two not previously tested: `gpt-5.4` and `o3`. Results are in Table XI.

The results reveal that **L4 resistance is inversely correlated with task complexity**. Three of four models achieve 100% pass rate (30/30); `gpt-4o` achieves 80% (24/30, with H1 at 4/10). This contrasts sharply with the original L4 results, where *all* models achieve 0% pass rate on the simple Fibonacci task (T1) under example-only implicit delivery.

Methodological note on comparability. This experiment differs from the strict L4 ablation (§V-D) in two respects, not one: (1) *prior entrenchment depth*—hard tasks have more varied implementations; and (2) *I/O specification transparency*—hard-task examples expose the behavioral target unambiguously via input/output pairs (e.g., [3,1,4,2]→[4,3,2,1] makes "descending order" self-evident), while the original L4 embeds the inversion as an implicit semantic rule in code structure. These two confounds co-vary: the harder tasks that have shallower priors *also* have more behaviorally transparent I/O specifications. We cannot cleanly separate prior-depth effects from I/O-transparency effects in the current design. Readers should interpret the hard-task results as evidence for *at least one* of these explanatory factors, not necessarily prior depth alone.

Prior entrenchment depth as primary explanatory variable (hypothesis). We hypothesize that the key variable is *prior entrenchment depth*: how uniformly a single canonical implementation dominates the training distribution. Fibonacci recursion has an extremely deep, uniform Python prior—virtually all training examples use the same `if n <= 1: return n; return fib(n-1) + fib(n-2)` template. Merge sort, binary search, and BFS admit diverse implementation strategies (iterative vs. recursive, different data structures, varied edge-case handling), yielding shallower and more varied priors. When the prior is shallow, the model can override it by extracting the inverted semantics from worked examples; when the prior is deeply entrenched, the NTP distribution overwhelms in-context cues. We label this a hypothesis because the I/O-transparency confound cannot be ruled out with current data.

A complementary observation reinforces this framing: the hard-task prompts provide *explicit I/O specifications* (descending-sorted output, inverted return values, reversed traversal order), while the original L4 tasks embed an *implicit semantic rule* in code structure. Models readily comply with behavioral specifications for complex algorithms, but cannot extract the same inversion when it is embedded implicitly in prior-entrenched code patterns.

`gpt-5.4` and `o3` hard-task results, and E24 update: Both achieve perfect 30/30 on the hard-task battery (shallow-prior, I/O-transparent tasks). To directly test whether these models also overcome *deep* prior entrenchment, we ran E24: strict example-only Variant A on L4 Fibonacci for both models ($n = 10$ each). Result: `gpt-5.4` 0/10 (PPR=1.00) and `o3` 0/10 (PPR=1.00). Frontier models are defeated by the same Fibonacci task that defeats all seven core models. The open question from the hard-task section is now resolved: *pattern blindness on prior-entrenched tasks holds for current frontier models* as of 2026-03-25.

VI. Discussion

A. Effective LLM-resistant language design

Our results identify a principled recipe for LLM-resistant programming languages:

1) Token-level perturbation alone is insufficient (models partially resist, but at measured rates).
2) Syntactic restructuring is also insufficient (models adopt explicit structural rules).
3) Semantic inversion with explicit instruction provides capacity-dependent resistance.
4) Semantic inversion without explanation (L4) provides consistent resistance via pattern blindness.
5) **Task selection is critical**: tasks with deep recursive priors (e.g., Fibonacci) are universally resistant, while tasks with shallower or more varied idioms (e.g., is_sorted) remain extractable by mid-to-large models. Instructors should prefer tasks with strong, uniform Python prior entrenchment.
6) **Prior entrenchment depth governs resistance**: L4 resistance holds only for tasks with deep, uniform Python priors (e.g., Fibonacci recursion template). Complex tasks with varied implementations (merge sort, BFS, binary search) allow capable models (`gpt-5.4`, `o3`, `o4-mini`) to extract inverted semantics from I/O examples—but this is partly confounded by I/O specification transparency (see §V-N). Critically, `gpt-5.4` and `o3` tested directly on strict L4 Fibonacci (E24) both fail: 0/10 (PPR=1.00)—confirming that even frontier models are defeated by prior-entrenched tasks, not just less capable models. E31 (§V-J) provides direct logprobs evidence: both fib and merge_sort have PES$\geq$0.9996, yet only fib resists L4—ruling out a simple prior-magnitude explanation and supporting I/O transparency as the critical moderator. Instructors should prefer tasks with strong, uniform Python prior entrenchment.
7) **Cross-task contamination breaks isolation** (E32, §V-K): L4 resistance is a prompt-level property, not just a task property. Showing inverted code examples from *other* tasks in the same prompt enables models to crack Fibonacci (`gpt-4.1-mini`: 10/10, up from 0/50 in isolation). Instructors must ensure prompt-level isolation of L4 assignments.
8) **Pattern mimicry dominates rule following** (E34, §V-M): showing an inverted code example (0% verbal annotation) produces higher pass rates than full explicit rules (100% annotation) across all tested models. This inverts the naive expectation and has direct implications for confusion language design: resistance depends on what the prompt *shows*, not what it *says*.

The most practical target for educational use is L4: retain Python surface syntax, invert conditional semantics, and provide only worked examples rather than explicit rules—applied to tasks with deep Python priors (e.g., Fibonacci). **Warning:** harder algorithmic tasks (merge sort, BFS, binary search) do *not* reliably provide L4 resistance because they have both shallower priors *and* more transparent I/O specifications; L4 works for prior-entrenched, pattern-blind tasks.

Two complementary resistance strategies. Our results identify two distinct, validated pathways to NTP-resistant assessment design. (1) *L4 prior exploitation*: retain Python syntax, invert semantics implicitly via worked examples, and choose tasks with deep Python priors—resistance arises because the model's NTP prior cannot be overridden by in-context example patterns. (2) *I/O opacity* (E35): design tasks with opaque input/output mappings where no training distribution provides a shortcut—resistance arises because models cannot reverse-engineer the function from I/O pairs alone, regardless of prior strength. These strategies target different failure modes (pattern blindness vs. generalization failure) and can be combined: an I/O-opaque task with an additional L4-style inversion would engage both failure modes simultaneously. Educators should select the strategy (or combination) appropriate to their task set and threat model.

L4 experimental conditions. The paper reports L4 findings across three distinct experimental conditions with different evidentiary strength; see §V-E for the full protocol note. The core claim—that L4 pattern blindness defeats current models—rests on the strict example-only ablation (Table IV): 9 models, 0 pass rate, no partial annotation. The partially-annotated multi-task (Table V) and hard-task extension (Table XI) provide enriching context about factors that modulate resistance but are not directly comparable to the headline ablation results.

B. Implications for programming education

Why human students retain an advantage. The practical value of L4 rests on a fundamental asymmetry that does not require experiment to accept. An instructor who tells students *"in this assignment, if-blocks execute when the condition is false—here are three worked examples"* is providing a rule that a human learner can read, verify against the examples, and apply step-by-step. The *same* prompt given to an LLM triggers a different mechanism: the model's next-token prediction (NTP) prior—built from billions of standard Python tokens—systematically pulls output back toward `if n <= 1: return n`, regardless of the stated rule. This is the core mechanism we measure: not that LLMs are "less intelligent," but that their generation process is governed by distributional priors that cannot be fully overridden by in-context instructions for sufficiently entrenched patterns. A student *reasons* about a rule; an LLM *samples* from a prior.

Our findings partially address the concern of LLM-solved homework. L4-style assignments provide meaningful resistance for tasks with deep Python priors (e.g., Fibonacci recursion): all non-reasoning models fail T1 (0/10 to 0/20 pass rates), though via different failure modes: prior dominance (PPR=1.0) for most models, partial template fragmentation (PPR=0.20) for `gpt-4.1-nano`; o4-mini achieves 8/15 (Wilson 95% CI: [33%, 77%]) only under the partially-annotated multi-task prompt (not under the strict example-only condition); the wide CI reflects the small $n = 15$ and should be interpreted as a noisy partial success rather than a definitive estimate. However,

task selection matters: T2 (is_sorted) is solved by capable models (`gpt-4o`, `gpt-4.1-mini` at 100%), indicating that simpler idioms with shallow priors may remain vulnerable.

Students who understand the L4 rules (from lecture instruction) retain an advantage over LLM-assisted submissions: current models universally fail L4 Fibonacci (T1), while a human who reads the rule and examples can apply it correctly. The human–LLM performance gap follows directly from the mechanics of NTP-based generation and does not require empirical validation to accept as a design principle: a model that samples from a distribution over Python training tokens cannot override that prior through in-context instruction alone, while a human who reasons about a stated rule can. A formal human pilot study would add quantitative precision to this asymmetry but is not required to establish the core resistance claim.

C. Non-monotonic capacity and the gpt-4.1 vs. gpt-4.1-mini anomaly

The most surprising within-family finding is the T2 capacity inversion: `gpt-4.1-mini` passes 20/20 on `is_sorted` (PPR=0.0), while the nominally "larger" `gpt-4.1` achieves only 1/50 (Fisher's exact $p < 0.0001$). We offer two non-exclusive hypotheses (stated as hypotheses, not findings): (a) *RLHF over-alignment*: the `gpt-4.1` variant may have been trained with stronger code-generation RLHF reward shaping, causing deeper entrenchment of canonical Python patterns and making prior override more difficult even when pattern cues are present; (b) *instruction-tuning regime*: the "mini" variant may have seen a higher proportion of rule-following fine-tuning data (relative to raw code data), yielding stronger in-context instruction compliance at the cost of raw generation fluency.

E25 diagnostic (L3 T2 for gpt-4.1): To discriminate between a general inability to invert `is_sorted` semantics versus an L4-specific pattern extraction failure, we ran $n = 10$ L3 T2 evaluations for `gpt-4.1` (explicit rule, same task). Result: `gpt-4.1` passes 10/10 (PPR=0.0). The model correctly applies the L3 semantic inversion rule for `is_sorted` when given an explicit instruction—the exact failure mode it exhibits on L4 T2 (1/50) is absent under L3. This result directly supports hypothesis (a): the RLHF over-alignment explanation. `gpt-4.1`'s T2 failure is not due to difficulty understanding inverted `is_sorted` semantics; it is specifically a *pattern extraction failure under implicit delivery*—the model cannot extract the inversion rule from worked examples, even though it can follow it from an explicit rule statement. Both explanations predict that scale alone does not determine pattern extraction ability—fine-tuning decisions interact with model size in complex ways. We flag this as a priority for mechanistic interpretability follow-up.

D. Open-weight model failure modes

The open-weight results reveal two distinct failure modes absent from the OpenAI family (see §V-H for full data and CIs). `Llama-3.3-70B` exhibits *context-sensitive prior reversion* (introduced in §V-H): variant-level PPR varies from 0.00 on in-domain examples to 1.00 on domain-shifted examples, suggesting the model's prior engagement is gated by contextual similarity rather than rule-following capacity. The implication for assessment design is direct: any LLM with context-sensitive prior reversion would require *variant diversity* to be reliably fooled—a single example set may inadvertently select a variant that triggers prior reversion recovery. `Qwen3-32B` exhibits *general semantic meta-rule failure*: the model fails even under L3 (explicit rule), with 0/5 pass and PPR=0.80 (E26 pilot), *confirmed at* $n = 20$ (E27: 0/20 pass, PPR=0.40, Wilson 95% CI [0%, 17%]). The lower PPR at $n = 20$ compared to E26 reflects that many failures are *confused partial instruction-following*—internally inconsistent outputs rather than clean Python prior returns—suggesting a richer failure mode than simple prior dominance. This contrasts sharply with all tested OpenAI models, which pass L3 consistently—including `gpt-4.1` (10/10, E25)—but fail L4. The Qwen result suggests that L3/L4 failure taxonomy is architecture-dependent: some models exhibit L4-specific pattern blindness (rule-following capability present, pattern extraction absent), while others exhibit L3-level instruction-following failure for semantic inversions (a deeper deficit). Educators using open-weight models should test L3 compliance before assuming that L4 resistance is the binding constraint.

Assessment scope for L3-failing models. For models that fail L3 (e.g., `Qwen3-32B`), L4 resistance is vacuous from an assessment design perspective: the model is already blocked at the more fundamental level of explicit-rule following, not by the L4 implicit-delivery challenge. L4-style assessment design is most meaningful when applied to models that pass L3 (instruction-following capable) but fail L4 (pattern extraction absent)—i.e., models where the failure is specifically the inability to extract inverted semantics from examples, not a general deficit. Instructors designing L4-resistant assessments should first verify L3 compliance for any LLM they wish to target, to confirm that the assessment challenge is the intended one.

E. Limits

While the multi-task experiment (Section V) covers three tasks and seven models (345 total runs, de-duplicating scale-up runs), the task set is limited. The task-dependent pass rate for `gpt-4o` on T2 (is_sorted: 10/10 pass vs. T1 fib: 0/10) confirms that L4 resistance is not unconditional—tasks with shallow or varied Python idioms remain extractable by capable models. The PPR metric for L3/L4 is task-specific and requires per-task oracle semantics. Human baseline experiments (Exp-5) are needed to confirm that L4 does not also confuse human learners.

The *operational substitution* failure mode is confirmed as dominant for T3 at $n = 50$: 47/50 `gpt-4o` runs (94%; Wilson 95% CI: [83%, 99%]). Our structural judge detects known arithmetic patterns (e.g., `count -= 1; return -count`), achieving high recall on observed workarounds (47/50 detected). Novel arithmetic transformations not matching these patterns may evade detection; an execution-based judge would provide

exhaustive coverage of this failure mode. The scale-up results establish it as real and frequent, not a one-off artifact.

The L4 ablation ($n = 50$ per model, 350 total runs) strengthens the pattern blindness finding: zero pass rate holds across all five prompt variants and all seven models including the reasoning model `o4-mini`. One important caveat: `o4-mini` was evaluated at temperature=1 (an API constraint), while all other models used temperature=0 (deterministic). Cross-model comparisons involving `o4-mini` should therefore be interpreted with caution, as sampling variance at temperature=1 may contribute to PPR variance across variants.

The high PPR variance across variants (e.g., `gpt-4.1-mini` variant A PPR=0.10 vs. variant B PPR=0.00 vs. variant E PPR=1.00) warrants caution—some configurations may trigger different failure sub-modes that are not yet fully characterized. Notably, `gpt-4.1-mini` shows non-monotonic variant behavior: variant A (2 examples, in-domain) achieves PPR=0.10 while variant B (3 examples, in-domain) achieves PPR=0.00—lower prior dominance with more examples. We hypothesize that example count and domain jointly modulate the model's extraction capability: additional in-domain examples may partially suppress the pure Python prior without enabling full inversion (yielding fail-other rather than fail-prior), while domain shift (variants D/E) reinvigorates the prior by reducing contextual alignment. This interaction is consistent with observed data but cannot be confirmed at current $n = 10$ per variant. A parallel run of `gpt-5.4-mini` completed $n = 42$ valid evaluations (8 HTTP errors in variant E); all valid runs show PPR=1.0 and pass=0, consistent with full prior dominance across all variants. The overall PPR of 1.00 for `gpt-5.4-mini` represents the strongest prior dominance among all tested models.

Notably, an earlier preliminary run of `o4-mini` at $n = 20$ yielded 1/20 pass; at $n = 50$, this result does not replicate (0/50), indicating that the earlier observation was likely a statistical artifact and not evidence of systematic reasoning-based inversion.

Corrected `o4-mini` result: partial annotation enables reasoning-model extraction. An earlier multi-task run of `o4-mini` used `max_completion_tokens=400`, which truncated reasoning-model responses before the code block was emitted. After correcting to `max_completion_tokens=5000`, `o4-mini` achieves 8/15 on T1 (partially-annotated prompt). This revision has two implications: (a) the 0/45 result from the strict example-only ablation (Table IV) is unaffected and confirms pattern blindness under pure implicit delivery; (b) under partial annotation (inline comments that partially articulate the rule), `o4-mini`'s reasoning capability provides a meaningful advantage that dense models cannot replicate. We report both results transparently. This suggests that the boundary between L3 (explicit rule) and L4 (implicit) is not binary—degrees of partial annotation interact with model reasoning depth.

F. Next steps

- Exp-4 (partial): Multi-task results acquired for T2/T3; broader task coverage (nested control-flow, stateful logic) remains for future work.
- Exp-5: Human study comparing student performance vs. LLM performance on L4 tasks; critical for validating the educational claim.
- Operational substitution: an execution-based judge would improve detection precision beyond the current structural heuristic, particularly for novel arithmetic workarounds beyond the `count -= 1` pattern.
- Annotation density gradient: E28 established non-monotonic density for `o4-mini`; E34 (§V-M) extends to `gpt-4o` and `gpt-4.1-mini`, confirming the pattern generalizes across architectures. Follow-up: test `o3` and other reasoning models; vary prior entrenchment depth to determine whether the saturation threshold shifts with task difficulty.
- Prior-depth vs. I/O-transparency confound: E31 (§V-J) shows both fib and merge_sort have PES$\geq$0.9996 yet diverge on L4 resistance, supporting I/O transparency as the critical moderator. E33's confounded design has been addressed by E35 (§V-L), which shows that opaque I/O pairs alone (no algorithm description) produce universal resistance (0/30, 0%) independent of prior strength. Together, E31+E35 support a two-factor model of resistance: prior entrenchment depth and I/O transparency are both sufficient (but distinct) sources of model failure. E35 establishes I/O opacity as sufficient for resistance on a separate (opaque_counter) task; whether removing I/O transparency from the hard tasks (H1–H3) would recover resistance remains untested and is the most direct experiment to resolve the hard-task confound.
- Cross-task contamination mitigation: E32 (§V-K) shows that cross-task inverted examples break L4 isolation. Future work should test whether prompt-level firewalls (e.g., system-message separation of examples from target task) can restore resistance.

VII. Conclusion

We presented a four-level taxonomy (L1–L4) for designing LLM-resistant programming languages, grounded in the NTP prior structure of language models.

Our central finding is that resistance does not come from surface-level token changes or structural novelty—both of which LLMs can adapt to with explicit instruction. Robust resistance requires semantic inversion: the same Python-like surface, but inverted conditional semantics. When semantic inversion is delivered implicitly (via examples only), all tested models exhibit pattern blindness and default to Python priors.

We identify three distinct failure mechanisms: *prior dominance*, where the prior overrides the stated rule; *pattern blindness*, where the model fails to extract the semantic rule from examples; and *operational substitution*, where the model achieves correct output via arithmetic workaround rather than adopting the semantic inversion rule. These mechanisms have different implications for language design: prior dominance

S0 — Baseline system prompt (L1)

"Solve the given prompt exactly. Return ONLY code (no explanation, no markdown). Follow alias mapping strictly."

S1 — Context-pack system prompt (L1, ablation)

"Solve the given prompt exactly. Return ONLY code. Follow alias mapping strictly."

[Transformed language contract]
Use transformed aliases as listed in the prompt. Do NOT output original Python keywords for transformed entries. Output only executable code in the transformed language.

Fig. 5: L1 system prompts. S0 is the compact baseline; S1 adds an explicit alias contract to test whether contextual reinforcement reduces prior dominance.

L1 user prompt template

[Language Description]
Python grammar remains unchanged.
Keyword aliases are user-configurable strings.
At runtime, detected aliases are normalized back to Python.

Current mapping summary:
`if` → {alias} `in` → {alias} `def` → {alias}
`for` → {alias} `return` → {alias} `print` → {alias}

[Task]
1) Define function `fib(n)`.
2) If `n <= 1`, return `n`.
3) Otherwise return `fib(n-1) + fib(n-2)`.
4) Loop `i = 0..6`, print `(i, fib(i))`.

Fig. 6: L1 user prompt template. Alias values are varied across v1–v120, using tokens of differing length, semantic distance, and collision potential. The task is fixed as Fibonacci to control task difficulty.

can be partially mitigated by explicit context packs; pattern blindness requires implicit delivery; operational substitution may require execution-based evaluation to detect.

These findings provide a principled basis for designing programming course assessments that remain meaningful in the LLM era, where LLMs will fail on L4-style problems even as they trivially solve standard Python assignments.

VIII. Prompt and Evaluation Protocol

A. L1: Token Alias Substitution

Alias sampling strategy. Representative mappings from three prompt versions:

Token	v1 alias	v35 alias	v90 alias
`if`	`return`	anti-prior phrase	long natural phrase
`def`	`if`	`def`	`return`
`for`	long phrase	`xk`	long phrase
`return`	long phrase	long phrase	`out`
`print`	`do`	—	`qn`

L2 grammar specification (complete)

This is NOT Python. All Python keywords are FORBIDDEN.

`:define name [args] ->` (replaces `def name(args):`)
`<= value` (replaces `return value`)
`value ?cond: res ;else: alt ;end` (replaces `if`)
`loop var over iterable {body}` (replaces `for`)
`;;` closes function block

Fig. 7: L2 grammar specification as given to the model. Models must produce code using exclusively these constructs. Result: all models fully adopted the syntax (SIR=0.0), showing LLMs readily follow explicit structural rules.

L3 system prompt — explicit semantic rule

*"Follow the L3 language rules exactly. If-blocks run when condition is **FALSE**."*

L3 user prompt (complete)

L3 rule: Syntax is identical to Python, but:
RULE: Indented block executes when condition is **FALSE**.
RULE: Else-block executes when condition is **TRUE**.

Example:

```
if x > 5:
    print("small")  # runs when NOT x > 5
else:
    print("large")  # runs when x > 5
```

Task: Write `fib(n)` in L3. Intended logic: if $n \leq 1$ return n; else return `fib(n-1)+fib(n-2)`.
Think carefully: your if-block runs when condition is FALSE.

Fig. 8: L3 prompts. The rule is stated explicitly in the system message. Weaker models (gpt-4o-mini) still fail (PPR=1.0), while stronger models succeed—revealing a capability threshold for semantic meta-rule following.

B. L2: Syntax Inversion

C. L3: Semantic Inversion (Explicit)

D. L4: Semantic Inversion (Implicit)

E. Evaluation Metrics

- **KLR** (Keyword Leakage Rate): fraction of alias slots where original Python keyword was used. ACR $= 1 -$ KLR is a derived quantity.
- **PSS** (Partial Structural Score): weighted sub-component credit (alias: 0.4, parse: 0.25, skeleton: 0.25, format: 0.1).
- **PPR** (Python Prior Rate, L3/L4): fraction of outputs using Python-style condition without inversion.

F. Judge Protocol (L1)

1) Alias compliance: check that alias is present and original token is absent.

L4 user prompt (example-only delivery, no explicit rule)

Here are some examples in our programming language:

Example 1 — max function:

```
def max_val(a, b):
    if a > b:
        return b   % smaller when a > b
    return a
```

Output: `max_val(3,5)=3`, `max_val(7,2)=2`

Example 2 — even check:

```
def is_even(n):
    if n % 2 == 0:
        return False
    return True
```

Output: `is_even(4)=False`, `is_even(3)=True`

Task: Write `fib(n)` in this language. Target: `fib(0)=0`, `fib(1)=1`, `fib(n)=fib(n-1)+fib(n-2)`. Print `fib(0)` through `fib(6)`.

Fig. 9: L4 prompt (no system message). The semantic inversion rule is not stated; models must infer it from worked examples. Even gpt-4o, which succeeds at L3, defaults to canonical Python fib—exhibiting *pattern blindness*.

2) Normalization: replace aliases → Python keywords via word-boundary regex.
3) Parse check: `ast.parse()` on normalized code.
4) Skeleton check: verify `fib`, base case, recursion, and print loop.

G. Judge Protocol (L3/L4)

- **Python-prior**: model writes `if n <= 1: return n` (non-inverted) ⇒ PPR=1 (fail).
- **L3-correct**: model writes `if n > 1: return n` or equivalent ⇒ PPR=0 (pass).

IX. Complete Experiment Log

A. All experiments (chronological)

1) **E1–E4: L1 baseline, $n = 120$, 4 models (2026-03-23).** Models: gpt-4o-mini, gpt-4.1-mini, gpt-4.1-nano, gpt-5.4-mini. Files: `batch-v1-120.`*`model`*`.json`
2) **E5: L1 context-pack ablation, $n = 120$, gpt-4o-mini (2026-03-23).** Setting: CONTEXT_PACK=1. KLR: $0.423 \rightarrow 0.260$ (39% reduction from E1).
3) **E6: L1 baseline, $n = 120$, gpt-4o (2026-03-24).** KLR=0.366, highest structural compliance.
4) **Partial judge applied (2026-03-24).** Script: `partial_judge.py` applied to all L1 batches.
5) **E7: L2 syntax inversion, $n = 5$, gpt-4o-mini (2026-03-24).** Explicit grammar rules provided in system prompt. SIR=0.0; model adopted L2 syntax fully.
6) **E8: L3 semantic inversion, $n = 5$, gpt-4o-mini (2026-03-24).** Explicit inversion rule in system prompt. PPR=1.0: prior dominance, complete failure.
7) **E9–E10: L3 semantic inversion, $n = 10$, gpt-4o and gpt-4.1-mini (2026-03-24).** Both models: PPR=0.0 (successfully followed rule).
8) **E11: L4 implicit semantic, $n = 5$, gpt-4o (2026-03-24).** No system prompt; only 2 worked examples shown. Pattern blindness observed: model ignored inversion cue.
9) **E12: L1 context-pack, $n = 20$, gpt-4.1 family + o4-mini (2026-03-25).** Models: gpt-4.1 (KLR=0.295), gpt-4.1-mini (KLR=0.210), gpt-4.1-nano (KLR=0.378), o4-mini (KLR=0.265). Files: `l1-gpt41-contextpack-2026-03-25.json`, etc. Note: o4-mini requires `temperature=1` (API constraint).
10) **E13: L4 ablation, $n = 50$, 5 models (2026-03-25).** Models: gpt-4o, gpt-4o-mini, gpt-4.1-mini, gpt-5.4-mini; later gpt-4.1 and gpt-4.1-nano scaled to $n = 50$. Result: 0/50 pass for all five variants × all models. File: `L4-ablation-n50.`*`model`*`.2026-03-25.json`
11) **E14: L4 ablation, $n = 50$, o4-mini (2026-03-25).** Reasoning model; confirms 0/50 pass at scale (1/20 preliminary was a statistical artifact). File: `L4-ablation-n50.o4mini.2026-03-25.json`
12) **E15: L4 multi-task, $n = 10$–20 per task, 7 models (2026-03-25).** Tasks: T1 (fibonacci), T2 (is_sorted), T3 (count_vowels). Models: gpt-4o, gpt-4o-mini, gpt-4.1, gpt-4.1-mini, gpt-4.1-nano, gpt-5.4-mini, o4-mini (255 total runs). gpt-5.4-mini: 0/30 (T1 PPR=1.0, T2 PPR=1.0, T3 operational substitution). Files: `L4-multitask.`*`model`*`.2026-03-25.json`
13) **E16: L4 multi-task scale-up, $n = 50$, gpt-4o T3 and gpt-4.1 T2 (2026-03-25).** Rationale: preliminary $n = 10$ results showed 1 pass each; scaled to $n = 50$ to rule out statistical artifact. gpt-4o T3 (`count_vowels`): 0/50 pass; operational substitution in 47/50 runs (94%, Wilson 95% CI: [83%, 99%]). Confirms preliminary pass was artifact. gpt-4.1 T2 (`is_sorted`): 1/50 pass (Wilson 95% CI: [0.0%, 10.5%]); Fisher's exact vs. gpt-4.1-mini 20/20: $p < 0.0001$. Total multi-task runs (E15+E16+E19): 345 (255 initial + 80 scale-up `gpt-4o` T3 / `gpt-4.1` T2 + 10 `o4-mini` T2/T3 extension). Files: `L4-scaleup-*.n50.json`. Script: `run_targeted_scaleup.py`
14) **E17: L2/L3 expansion, $n = 20$ each, 3 models (2026-03-25).** Models: gpt-4o, gpt-4.1-mini, gpt-4.1-nano. Rationale: prior L2/L3 data was pilot-level ($n = 5$–10). L2 results: all three models 20/20 pass (SIR=0.0); confirms full syntax adoption at statistically meaningful n. L3 results: all three models 20/20 pass (PPR=0.0); `gpt-4.1-nano` 20/20 is notable—even the smallest model fully follows explicit rule. Key contrast: `gpt-4.1-nano` L3 20/20 vs. L4 0/50 (PPR=1.0); explicit vs. implicit delivery. Files: `l{2,3}-*-n20-*.json` (6 files). Scripts: `run_l{2,3}_eval.py`
15) **E18: Open-weight L1 + L4, $n = 20$, Llama-3.3-70B and Qwen3-32B via Groq (2026-03-25).** Models accessed via Groq API (User-Agent: "OpenAI/Python"). L1 ctx-pack:

Llama-3.3-70B KLR=0.23 (comparable to best OpenAI); Qwen3-32B KLR=1.00 (worst overall). L4 ablation Variant A pilot: both models 0/20 pass (Llama PPR=0.05, Qwen PPR=1.00); confirms cross-family pattern blindness. Full 5-variant ablation ($n = 50$ attempted): Qwen3-32B 50/50 valid, 0/50 pass, PPR=1.00; Llama-3.3-70B 25/50 valid (25 HTTP errors), 0/25 pass, per-variant PPR: A=0.00, B=0.00, C=0.60, D=1.00, E=1.00 (overall PPR=0.52). L1 family variance (0.23–1.00) exceeds gpt-4.x intra-family variance (0.21–0.69). Files: `l{1,4}-*-groq-*.json`, `L4-ablation-n50.groq.llama3370bversatile.2026-03-25.json`

16) **E19: o4-mini multi-task T2/T3 extension, $n = 20$ each (2026-03-25).** Rationale: E15 only ran T1 with corrected token budget; T2/T3 needed extension to $n = 20$. T2 (`is_sorted`): 0/20 pass, PPR=0.30; T3 (`count_vowels`): 0/20 pass, PPR=0.00. Confirms reasoning advantage is annotation-contingent (T1 partial annotation) and does not generalize to T2/T3. Total o4-mini multi-task runs after E15+E19: 15+20+20=55. File: `l4-multitask-o4mini-T2T3-n20.json`

17) **E20: CoT ablation, $n = 20$ per condition, 2 models (2026-03-25).** Prompt variant: L4 Variant A + "think step by step about the pattern first" prefix. gpt-4o: no-CoT 0/20 (PPR=1.00) vs. CoT 0/20 (PPR=0.15, `mentions_inversion`=100%). gpt-4.1-mini: no-CoT 0/20 (PPR=0.40) vs. CoT 0/20 (PPR=0.00, `mentions_inversion`=100%). Key finding: reasoning–generation dissociation; models articulate inversion rule in CoT but still generate standard Python prior. Files: `l4-cot-ablation-*.json`. Script: `run_l4_cot_ablation.py`

18) **E22: CoT ablation extension, $n = 20$ per condition, 2 additional models (2026-03-25).** Extended E20 to `gpt-4.1-nano` and `o4-mini` (160 total CoT runs across 4 models). `gpt-4.1-nano`: no-CoT 0/20 (PPR=1.00) vs. CoT 0/20 (PPR=0.20, `mentions_inversion`=100%). `o4-mini`: no-CoT 0/20 (PPR=0.95) vs. CoT 1/20 (PPR=0.65, `mentions_inversion`=35%). Key finding: reasoning–generation dissociation confirmed across all 4 model families (dense + reasoning). CoT PPR reduction is largest in nano (1.00→0.20) but pass rate remains 0/20 for 3 of 4 models. Files: `l4-cot-ablation-gpt41nano-*.json`, `l4-cot-ablation-o4mini-*.json`. Script: `run_l4_cot_ablation.py` (`EVAL_MODEL=gpt-4.1-nano/o4-mini`)

19) **E23: L4 hard-task battery, $n = 10$ per cell, 4 models, 3 tasks (2026-03-25).** Hypothesis: harder algorithmic tasks amplify pattern blindness. Tasks: H1 (merge_sort, descending), H2 (binary_search, inverted return), H3 (BFS, reversed traversal). Models: `gpt-4o`, `gpt-5.4`, `o3`, `o4-mini` (120 total runs). Result: hypothesis REJECTED. `gpt-5.4`, `o3`, `o4-mini`: 30/30 (100%); `gpt-4o`: 24/30 (80%). Key finding: hard tasks provide *explicit I/O specifications* allowing direct behavioral compliance; simple L4 tasks embed an *implicit semantic rule* that models cannot extract. Pattern blindness is specific to implicit rule extraction, not task complexity. Prior entrenchment depth (not algorithmic difficulty) is the key predictor of resistance. Files: `l4-hard-tasks-*.2026-03-25.json`, `hard-vs-simple-summary-2026-03-25.md`.

20) **E21: Failure taxonomy, 419 failed L4 runs (2026-03-25).** Classified all failed L4 ablation + multi-task T1 runs across ablation (342 valid runs, all failures; `gpt-5.4-mini` $n_{\text{eff}} = 42$) and multi-task T1 (77 failed runs; 8 `o4-mini` T1 passes excluded from taxonomy). Type-I (exact Python prior): 83.1%; Type-II (partial attempt): 12.6%; Type-III (operational substitution): 4.3%. Type-I dominates in `gpt-4.1`/`gpt-4.1-nano` (100%); Type-II concentrates in `gpt-4.1-mini` (44%) and `gpt-4o-mini` (32%). File: `failure-taxonomy-2026-03-25.json` Script: `scripts/categorize_failures.py`

21) **E25: L3 T2 diagnostic for `gpt-4.1`, $n = 10$ (2026-03-25).** Motivation: `gpt-4.1` achieves only 1/50 on L4 T2 (`is_sorted`), despite L3 T2 being solvable for other models. E25 tests whether this failure is (a) L4-specific pattern extraction failure, or (b) general inability to invert `is_sorted` semantics. Condition: explicit L3 rule ("if-block runs when condition is FALSE"), task `is_sorted`, $n = 10$. Result: `gpt-4.1` passes 10/10 (PPR=0.0). Model correctly applies inverted conditional logic under explicit rule. Interpretation: supports RLHF over-alignment hypothesis—`gpt-4.1`'s T2 failure is L4-specific (pattern extraction from examples), not a general `is_sorted` inversion difficulty. File: `e25-l3-t2-gpt41-n10-2026-03-25.json` Script: `scripts/run_e25_l3_t2_gpt41.py`

22) **E26: L3 Fibonacci diagnostic for `Qwen3-32B`, $n = 5$ pilot (2026-03-25).** Motivation: `Qwen3-32B` fails L4 Fibonacci with PPR=1.00 across all variants. E26 tests whether this reflects (a) L4-specific pattern blindness (same as OpenAI models) or (b) general semantic meta-rule failure (fails explicit L3 rule too). Condition: explicit L3 rule ("if-block runs when condition is FALSE"), Fibonacci task, $n = 5$, via Groq API. Result: 0/5 pass, PPR=0.80. Model fails even when the inversion rule is stated explicitly. Interpretation: `Qwen3-32B` exhibits *general semantic meta-rule failure*. Pilot result; see E27 for full replication. File: `e26-qwen3-32b-l3-fibonacci.n5.2026-03-25.json` Script: `scripts/run_e26_qwen_l3.py`

23) **E27: L3 Fibonacci replication for `Qwen3-32B`, $n = 20$ (2026-03-25).** Motivation: E26 ($n = 5$) was preliminary; E27 replicates at $n = 20$ to confirm *general semantic meta-rule failure* is systematic, not a small-n artifact. Condition: identical to E26 (explicit L3 rule, Fibonacci task), $n = 20$ via Groq API. 5 HTTP 429 errors; $n_{\text{eff}} = 15$ for rate-limited subset; 20 total attempts. Result: 0/20 pass, PPR=0.40 (8/20 exact Python prior; 12/20 FAIL-other). Wilson 95% CI: [0%, 17%]. Interpretation:

Confirmed *general semantic meta-rule failure* is systematic at $n = 20$. Lower PPR than E26 reflects *confused partial instruction-following* (internally inconsistent outputs), not clean prior dominance. Architecture-level deficit, distinct from all OpenAI models. File: `e27-qwen3-32b-l3-fibonacci.n20.2026-03-25.json` Script: `scripts/run_e27_qwen_l3_n20.py`

24) **E24: Frontier model L4 Fibonacci ablation, $n = 10$, `gpt-5.4` and `o3` (2026-03-25).** Motivation: E23 only tested `gpt-5.4`/`o3` on hard tasks (I/O-transparent, shallow priors). C2 critique required testing same models on strict L4 Fibonacci to confirm core claim. Condition: Variant A (2ex-baseline), strict example-only, identical to Table IV. `gpt-5.4`: 0/10 pass, PPR=1.00 (complete prior dominance). `o3`: 0/10 pass, PPR=1.00 (complete prior dominance). Key finding: L4 pattern blindness on prior-entrenched Fibonacci holds for both frontier models as of 2026-03-25. Files: `e24-frontier-l4-fibonacci.gpt-54.n10.2026-03-25.json`, `e24-frontier-l4-fibonacci.o3.n10.2026-03-25.json`. Script: `scripts/run_e24_frontier_l4.py`

B. API and runtime settings

Setting	Value
Temperature	0 (all); 1 (o4-mini)
`max_compl_tokens`	default; 5000 (o4-mini)
Base URL	OpenAI `/v1`
Groq URL	Groq `openai/v1`
Timeout	120s
Prompt delay	0.4s
API key	Env variable
Retry	None
Python	3.12
`openai` pkg	1.75.0
`groq` pkg	0.22.0
`latexmk`	4.87

C. Repository structure for replication

Script / Path	Purpose
`run_gpt54_eval.py`	L1 batch runner
`run_l2_eval.py`	L2 batch runner
`run_l3_eval.py`	L3 batch runner
`run_l4_ablation.py`	L4 ablation runner
`run_l4_multitask.py`	L4 multi-task runner
`run_targeted_scaleup.py`	L4 targeted scale-up ($n = 50$)
`run_l4_cot_ablation.py`	CoT ablation runner
`categorize_failures.py`	Failure type taxonomy
`partial_judge.py`	Continuous metrics
`compare_eval_runs.py`	Run comparison
`docs/archive/*/`	Prompt banks (L1–L4)
`docs/research/results/`	All output JSONs
`language-levels/`	Level specs

X. Appendix: Prompt Samples, Taxonomy, and Config Snapshot

A. Prompt samples (readable form)

Representative prompt-map blocks from the versioned bank:

Sample A (v1)

- if → return
- def → if
- for → long natural phrase (renamed symbol)
- in → long natural phrase (strict compilation token)
- return → long natural phrase (keyword mapping token)
- print → do

Sample B (v35)

- in → mapped token phrase
- def → def
- for → xk
- return → long phrase preserving grammar
- if → anti-prior phrase

Sample C (v90)

- def → return
- for → phrase mapped back to reserved token
- in → exact mapping phrase
- return → out
- if → long natural phrase
- print → qn

B. Violation taxonomy summary (gpt-4o-mini, v1–v120)

From the context-pack comparison artifact, dominant violation categories are:

- Original keyword leakage (e.g., `for`, `if`, `return`, `in`, `def`)
- Parse failure after normalization (line-specific syntax failure)
- Missing alias emissions for required transformed tokens

Top counts in context-pack run include:

- `for` leaked: 52
- parse failed (line 1): 35
- `if` leaked: 33

C. PSS (Partial Structural Score) Supplementary Data

PSS is a heuristic composite metric (alias compliance 0.4, parse 0.25, skeleton 0.25, format 0.1) reported as a secondary diagnostic only. Weights are not validated; PSS is not used as a primary result. The table below provides reference values alongside KLR for L1 evaluation.

D. L4 Multi-Task Full PPR Values

Full Python Prior Rate (PPR) values for Table V. PPR = fraction of outputs matching canonical Python prior (lower = more L4-compliant). Pass rate = 0 for all cells except `gpt-4o` T2, `gpt-4.1-mini` T2, and `o4-mini` T1.

E. Execution config snapshot (non-secret)

Secret-bearing values (API keys) are excluded and loaded via runtime environment. Language specs are in `language-levels/` (repository).

TABLE XII: L1 KLR and PSS (supplementary). PSS is exploratory; KLR is the validated primary metric.

Model	Delivery	KLR↓	PSS↑
gpt-4.1-mini	ctx-pack	0.21	0.37
Llama-3.3-70B⋆	ctx-pack	0.23	0.34
gpt-4o-mini	ctx-pack	0.26	0.32
o4-mini	ctx-pack	0.27	0.41
gpt-4.1	ctx-pack	0.30	**0.43**
gpt-4o	baseline	0.37	0.40
gpt-4.1-nano	ctx-pack	0.38	0.30
gpt-4o-mini	baseline	0.42	0.27
gpt-4.1-mini	baseline	0.43	0.33
gpt-5.4-mini	baseline	0.67	0.25
gpt-4.1-nano	baseline	0.69	0.16
Qwen3-32B⋆	ctx-pack	1.00	0.00

⋆Open-weight via Groq API; $n = 20$, ctx-pack only.

TABLE XIII: L4 multi-task full PPR. † Partially-annotated prompts (not comparable to Table IV).

	T1: fibonacci			T2: is_sorted			T3: count_vowels		
Model	n	Pass	PPR	n	Pass	PPR	n	Pass	PPR
gpt-4o	10	0/10	1.00	10	10/10	0.00	50	0/50	0.00
gpt-4o-mini	10	0/10	0.00	10	0/10	1.00	10	0/10	0.00
gpt-4.1	10	0/10	1.00	50	1/50	0.42	10	0/10	0.00
gpt-4.1-mini	20	0/20	1.00	20	20/20	0.00	20	0/20	0.25
gpt-4.1-nano	10	0/10	0.20	10	0/10	1.00	10	0/10	0.00
gpt-5.4-mini	10	0/10	1.00	10	0/10	1.00	10	0/10	0.00
o4-mini	15	8/15	0.33	20	0/20	0.30	20	0/20	0.00

F. Experiment Log: E31–E34 (2026-03-25)

E31 (Prior Entrenchment Score). Logprobs API query at base-case decision point (`max_tokens=1`, `top_logprobs=5`, `temperature=0`). Models: `gpt-4o`, `gpt-4.1-mini`, `gpt-4o-mini`. Tasks: fib, merge_sort, is_sorted, count_vowels. 12 API calls total. `gpt-4o-mini` excluded from analysis (emits markdown backtick tokens rather than code tokens at this position). Results: `docs/research/results/e31-prior-entrenchment-lc3.5.json`.

E32 (Cross-Task Transfer). Few-shot prompts with inverted merge_sort and binary_search code, then fib request. Models: `gpt-4o`, `gpt-4.1-mini`, `o4-mini`. $n = 10$ per model (initial), temperature=0.7. Judge: regex check for inverted comparison operators ($>, \geq$) in base case. Initial `o4-mini` run yielded HTTP 400 (missing `max_completion_tokens`); re-run with corrected parameters yielded 1/10. `gpt-4o` subsequently scaled to $n = 30$ (two additional 10-run rounds, E32a/E32b) to tighten the confidence interval; combined result: 16/30 (Wilson 95% CI: [36%,70%]). Results: `docs/research/results/e32-cross-task-transfer.json`, `e32-gpt4o-scaleup.json`, `e32-gpt4o-scaleup2.json`.

E33 (I/O Transparency). Two tasks: IO-A (reverse_fib, transparent) and IO-B (opaque_counter, opaque). Models: `gpt-4o`, `gpt-4.1-mini`, `gpt-4o-mini`. $n = 10$ per task per model, temperature=0.7. Null result: 10/10 on both tasks for all models. Design flaw: IO-B prompt included explicit algorithm, making it procedurally transparent. Results: `docs/research/results/e33-io-transparency.json`.

E34 (Annotation Density Curve). Five density levels (0%, 25%, 50%, 75%, 100%) on L4 Fibonacci with inverted comparisons. Models: `gpt-4o`, `gpt-4.1-mini`, `o4-mini`. $n = 10$ per level per model (150 total). Temperature=0.7. Judge: regex check for inverted comparison in base case. Results: `docs/research/results/e34-annotation-density.json`.

TABLE XIV: Key runtime settings

Setting	Value
Models	gpt-4o, gpt-4o-mini, gpt-4.1, gpt-4.1-mini, gpt-4.1-nano, gpt-5.4-mini, o4-mini
L1 range	v1–v120 (baseline); v1–v20 (ctx-pack for 4.1/o4-mini)
CONTEXT_PACK	0/1
Temperature	0 (all models except o4-mini); 1 (o4-mini, API requirement)
Judge	deterministic

XI. Reproducibility and Artifact Availability

All experiments are committed to the repository. Representative paths:

- Results: `docs/research/results/`
- Language specs: `language-levels/`
- Evaluation scripts: `scripts/partial_judge.py`, `scripts/run_l2_eval.py`
- Paper: `paper/overleaf/`

References

[1] M. Chen *et al.*, "Evaluating large language models trained on code," *arXiv preprint arXiv:2107.03374*, 2021.

[2] J. Austin *et al.*, "Program synthesis with large language models," *arXiv preprint arXiv:2108.07732*, 2021.

[3] T. Zhuo *et al.*, "BigCodeBench: Benchmarking code generation with diverse function calls and complex instructions," arXiv:2406.15079, 2024.

[4] N. Jain *et al.*, "LiveCodeBench: Holistic and contamination-free evaluation of large language models for code," arXiv:2403.07974, 2024.

[5] Z. Feng *et al.*, "CodeBERT: A pre-trained model for programming and natural languages," in *EMNLP-Findings*, 2020.

[6] Y. Wang, W. Wang, S. Joty, and S. C. Hoi, "CodeT5: Identifier-aware unified pre-trained encoder-decoder models," in *EMNLP*, 2021.

[7] D. Fried *et al.*, "InCoder: A generative model for code infilling and synthesis," 2022.

[8] Z. Luo *et al.*, "WizardCoder: Empowering code large language models with evol-instruct," arXiv:2306.08568, 2024.

[9] D. Guo *et al.*, "DeepSeek-Coder: When the large language model meets programming," arXiv:2401.14196, 2024.

[10] Y. Liu *et al.*, "CodeIF: Benchmarking instruction-following in code generation," arXiv:2502.19166, 2025.

[11] W. Zhang *et al.*, "MultiCodeIF: Comprehensive evaluation of code instruction-following," 2025, preprint.

[12] E. Chen *et al.*, "The instruction hierarchy: Training LLMs to prioritize privileged instructions," arXiv:2404.13208, 2024.

[13] J. Xie *et al.*, "Adaptive chameleon or stubborn sloth: Revealing the behavior of LLMs in knowledge conflicts," ICLR 2024 (arXiv:2305.13300), 2024.

[14] S. Wu *et al.*, "Untangle the KNOT: Interweaving conflicting knowledge and reasoning skills in LLMs," arXiv:2403.08319, 2024.

[15] Z. Zhang *et al.*, "Knowledge conflicts for LLMs: A survey," EMNLP 2024 (arXiv:2403.08319), 2024.

[16] F. Shi *et al.*, "Large language models can be easily distracted by irrelevant context," ICML 2023 (arXiv:2302.00093), 2023.

[17] Z. Wang *et al.*, "A law of next-token prediction in large language models," arXiv:2408.13442, 2024.

[18] S. Min *et al.*, "Rethinking the role of demonstrations: What makes in-context learning work?" EMNLP 2022 (arXiv:2202.12837), 2022.
[19] Q. Dong *et al.*, "A survey on in-context learning," arXiv:2301.00234, 2023.
[20] A. Lampinen *et al.*, "A-Not-B errors in LLMs: Limitations of in-context learning," arXiv:2409.15454, 2024.
[21] Q. Zhao *et al.*, "LLMs struggle to perform counterfactual reasoning with parametric knowledge," arXiv:2410.06392, 2024.
[22] M. Turpin *et al.*, "Language models don't always say what they think: Unfaithful explanations in chain-of-thought prompting," NeurIPS 2023 (arXiv:2305.04388), 2023.
[23] K. Valmeekam *et al.*, "LLMs still can't plan; can LRMs? evaluation of OpenAI's o1," arXiv:2409.13373, 2024.
[24] K. Zhu *et al.*, "PromptBench: Evaluating robustness of large language models," arXiv:2306.04528, 2023.
[25] R. Pan *et al.*, "Attacking LLMs with code generation: Semantically-aligned adversarial perturbations," arXiv:2412.16359, 2024.
[26] S. Jha *et al.*, "Break the chain: Reasoning failures in LLMs via adversarial code prompting," arXiv:2506.06971, 2024.
[27] A. Khan *et al.*, "Defeating LLMs with adversarial perturbations in programming assignments," arXiv:2410.09318, 2024.
[28] J. Wu *et al.*, "A survey on LLM-based code generation for low-resource and domain-specific languages," arXiv:2410.03981, 2024.
[29] W. Chen *et al.*, "ConCodeEval: Evaluating LLMs for code constraint comprehension in DSLs," EMNLP 2024, 2024.
[30] A. Ni *et al.*, "L2CEval: Evaluating language-to-code generation capabilities," TACL 2024, 2024.
[31] Y. Dong *et al.*, "XGrammar: Flexible and efficient structured generation for LLMs," arXiv:2411.15100, 2024.
[32] J. Muendler *et al.*, "Type-constrained code generation with language models," PLDI 2025, 2025.
[33] C. Jimenez *et al.*, "SWE-bench: Can language models resolve real-world GitHub issues?" 2024.
[34] A. Dakhel *et al.*, "GitHub Copilot AI pair programmer: Asset or liability?" 2023.
[35] S. Lau and P. Guo, "From coding assistant to software developer: How novice programmers use LLMs," arXiv:2304.01984, 2023.
[36] J. Prather *et al.*, "It's complicated: Perspectives of novice programmers on GitHub Copilot," arXiv:2303.08733, 2023.
[37] M. Liffiton *et al.*, "CodeHelp: Using large language models with guardrails for programming classes," arXiv:2308.06921, 2023.
[38] E. Novak and B. McDanel, "Designing LLM-resistant programming assignments," SIGCSE TS 2025, 2025.
[39] M. Orenstrakh *et al.*, "Detecting LLM-generated code assignments using perplexity of code," arXiv:2307.09714, 2023.
[40] L. Ouyang *et al.*, "Training language models to follow instructions with human feedback," NeurIPS 2022 (arXiv:2203.02155), 2022.